DEEP LEARNING – WS 2019/20

LECTURE SCHEDULE

Date	Topics		Homework
25.11.	Math & ML Basics	Review of math basics ▪ linear algebra ▪ Probability ▪ information theory	
28.11.		Review of machine learning basics ▪ Supervised, unsupervised, reinforcement learning ▪ Empirical risk minimization ▪ Overfitting (05_ml:2-7) ▪ No free lunch theorem	Intro to Jupyter + tech stack Simple linear/logistic regression example in numpy ▪ (Batch) gradient descent manually, not AutoDiff ▪ Separation of training, validation & test sets ▪ Overfitting ▪ (L2) regularization **Due date: 16.12.**
16.12.	Perceptron and multi-layer perceptron	▪ Definition of perceptron and MLP ▪ Why deep? (01_intro:2-4, 06_mlp:17-20) ▪ Training deep neural networks ○ Optimization -> gradient descent (4.3, 04_numerical) ▪ Backpropagation (06_mlp) ▪ Stochastic gradient descent (5.9, 8.4+8.5) ▪ Vanilla SGD ▪ Momentum ▪ RMSProp ▪ Adam ○ Regularization ▪ Early stopping ▪ L1/L2 ▪ Dropout ○ Hyperparameter optimization: grid vs. random search	
19.12.			▪ MLP on relatively small nonlinear dataset ▪ Using PyTorch ▪ Explore different regularizers via random search ○ Early stopping ○ L2 ○ Dropout **Due date: 6.1.**

6.1.	Convolutional neural networks (CNN)	▪ Motivation and definition (9)
		  o weight sharing
		  o sparse connectivity
		  o growing receptive fields
		  o pooling for invariance
		  o striding, dilation, deconvolution
		  o padding
		▪ Architectures of classic and modern image processing CNNs: LeNet, AlexNet, VGG, ResNet, Inception, DenseNet, NAS
		▪ More on optimization
		  o Initialization
		  o Vanishing and exploding gradient problem
		  o Batch normalization as a simple fix for CNNs
		▪ Timeseries (1D) and video (3D)
			▪ CNN on (fashion) MNIST, CIFAR-10 or similar
			▪ Explore architecture choices like number of features, kernel sizes, pooling, striding, padding
			▪ Batch normalization
			Due date: 13.1.

9.1.			

13.1.	Transfer Learning & perceptual loss functions	▪ Pre-training on ImageNet
		  o DeCaf paper
		  o Object detection, segmentation
		  o Saliency
		▪ Deep features as perceptual loss functions
		  o Texture synthesis
		  o Style transfer
			Fine-tune or train classification layer of pre-trained ResNet on an image classification task with relatively little data available
			Due date: 20.1.

16.1.	Deep Image Synthesis	▪ Deep generative image models
		  o Autoencoders
		  o (Convolutional) Autoencoder
		  o VAE + beta VAE + spatial broadcast VAE
		▪ Autoregressive models
		  o Mixture models
		  o Spatial LSTM
		  o PixelCNN
		▪ Generative adversarial networks (GAN)	

20.1.		▪ Image-to-image models
		  o Domain translation networks
		  o Style transfer
		  o Superresolution
			▪ VAE on MNIST
			▪ Visualize classes in latent space
			▪ Visualize latent traversals
			Due date: 27.1.

23.1.	Recurrent neural networks (RNN) and text processing	▪ Basics/theory
		  o Unfolding computational graph (10.1)
		  o Simple RNNs (10.2)
		  o Bidirectional RNNs (10.3)
		  o Sequence-to-sequence models (10.4)
		  o Deep RNNs (10.5)
		▪ Recurrent architectures used in practice
		  o Gated recurrent unit (GRU)
		  o Long-short-term memory module (LSTM)
		  o Neural Turing Machine
		  o Seq2Seq: translation example
		▪ Deep word and text embeddings
		  o Word2Vec
		  o Transformer
		  o ELMo + BERT + GPT-2	

27.1.			▪ Text processing problem (topic classification, sentiment analysis)
			▪ Implement solutions
			  o based on LSTM
			  o using transfer learning using from BERT/GPT-2
			Due date: 3.2.

30.1.
3.2.
	Deep Reinforcement Learning	▪ Review Reinforcement Learning
		▪ (Deep) RL approaches
		  o Vanilla Policy Gradient
		  o Trust Region Policy Optimization
		  o Proximal Policy Optimization
		  o Deep Deterministic Policy Gradient (DDPG)
		  o Twin Delayed DDPG
		  o Soft Actor-Critic
		▪ Practical example: AlphaGo/AlphaZero	

6.2.	Backup/review		