

Deep Learning

Lecture 8: Transfer learning & perceptual losses

Alexander Ecker
Institut für Informatik, Uni Göttingen



<https://alexanderecker.wordpress.com>

– Credit: some of the slides based on Fei Fei Li, Justin Johnson and Serena Yeung's slides –

Where to put the softmax?

```
nn.Sequential(  
    nn.Linear(28*28, 28*28),  
    nn.ReLU(),  
    nn.Linear(28*28, 14*14),  
    nn.ReLU(),  
    nn.Linear(14*14, 10)) ← No softmax!  
...  
loss_fn = F.cross_entropy ← Includes softmax!  
...  
  
y_pred = model(x)  
correct = y == y_pred.argmax(dim=1)  
loss = loss_fn(y_pred, y)  
loss.backward()
```

Softmax is usually included in cross entropy for numerical reasons.

Consider what happens with float32 if you compute probabilities via softmax and then take logs for loss function.

Today

This is us



Standing on the shoulders of giants

Today's topics

1. Transfer learning
2. Neural style transfer

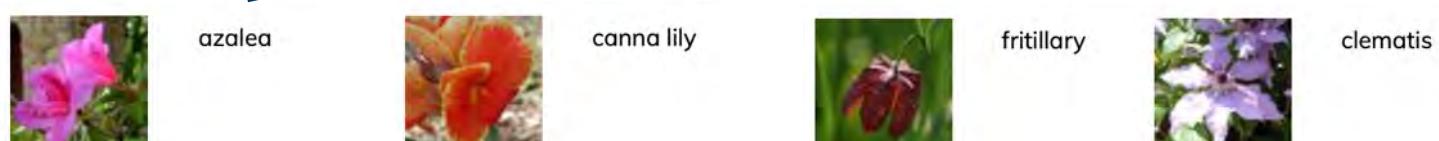
Transfer learning

-

Transfer learning: motivation

Illustrative example

Suppose you want to build an app to recognize flowers...



1000 images
100 classes

...

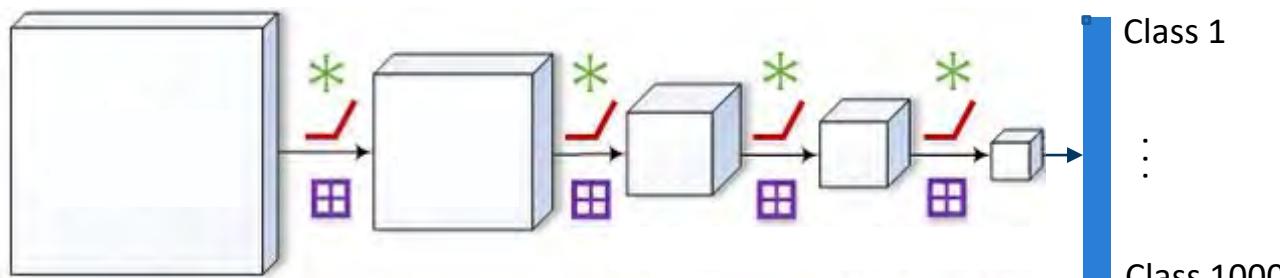
Training a convolutional net from scratch won't work well

Transfer learning: the idea

1

Pre-training
(massive data)

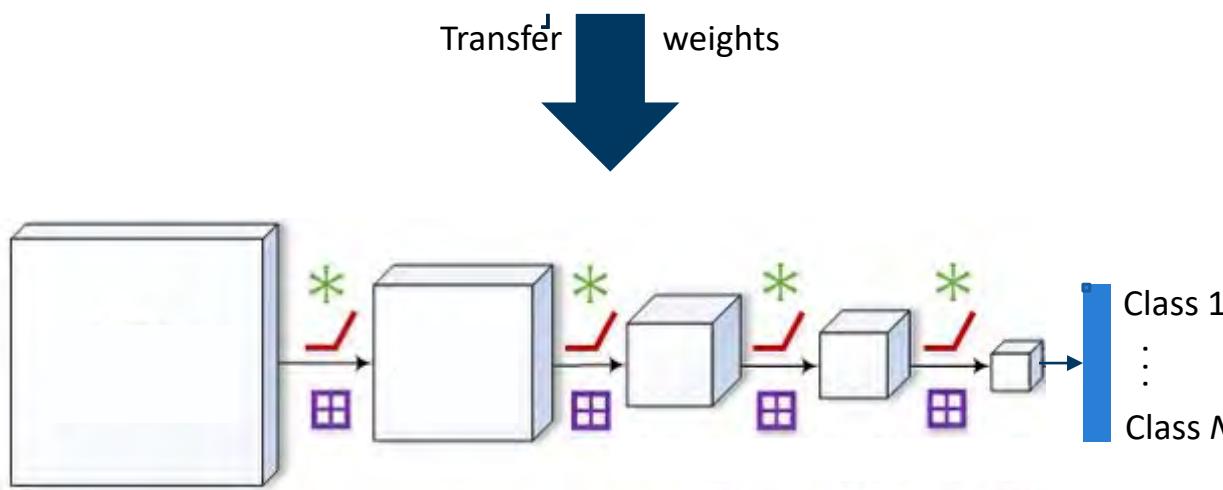
IMAGENET
1.2 million images
1000 object classes



IMAGENET

2

Fine-tuning
(much less data)



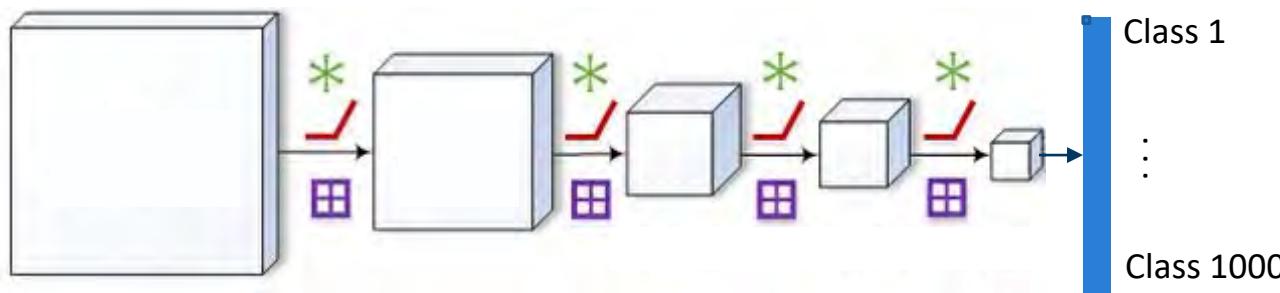
Recognize flowers

Transfer learning: the idea

1

Pre-training
(massive data)

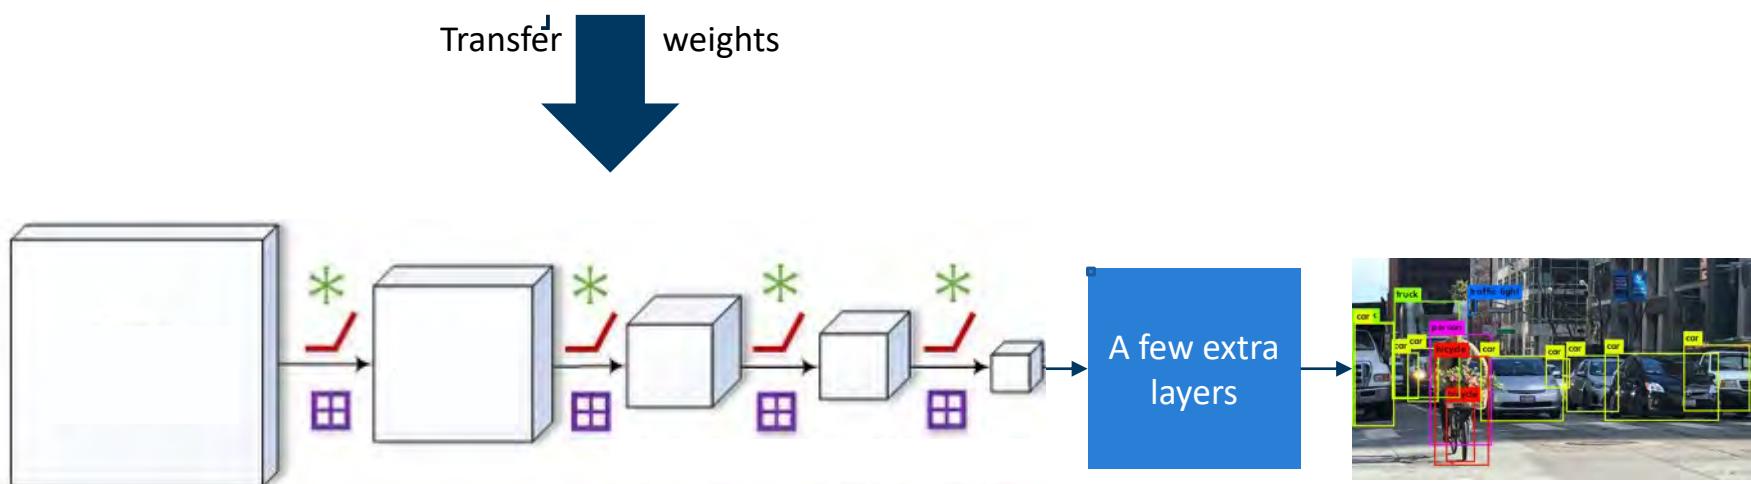
IMAGENET
1.2 million images
1000 object classes



IMAGENET

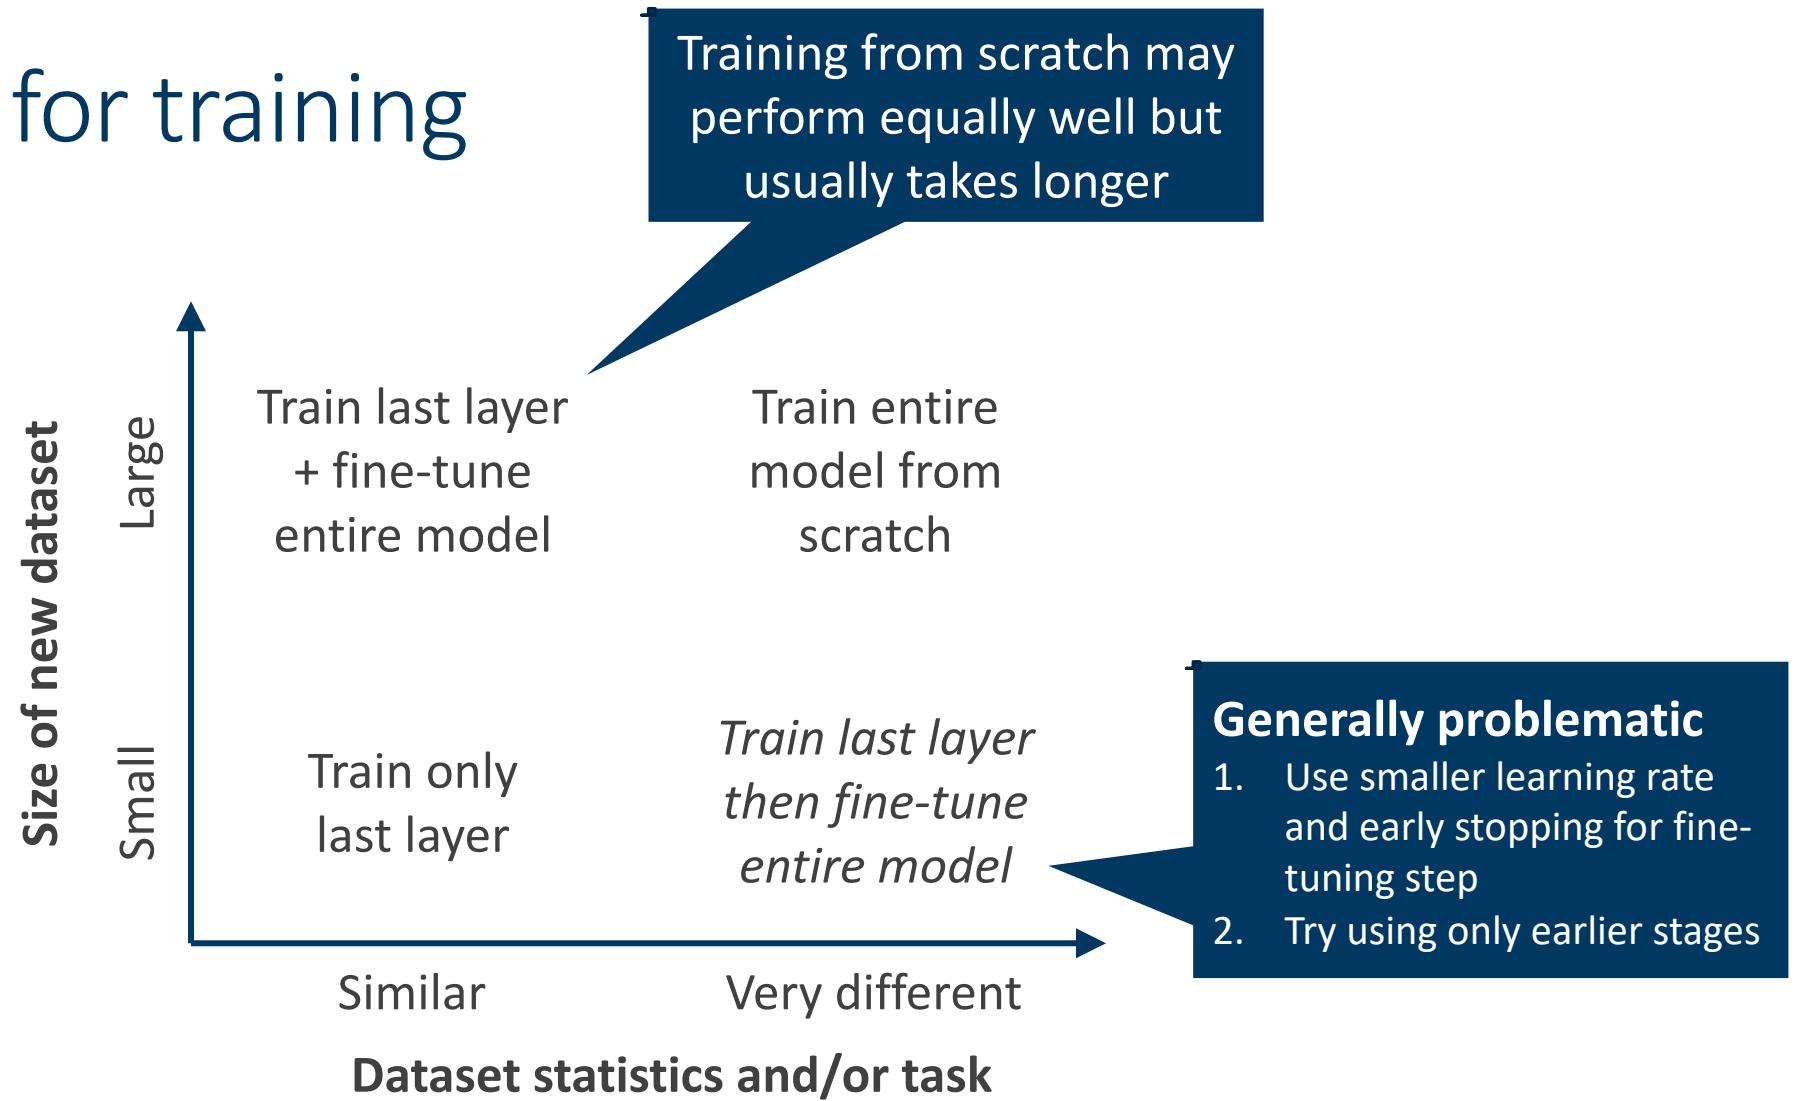
2

Fine-tuning
(much less data)



Object detection

Strategies for training



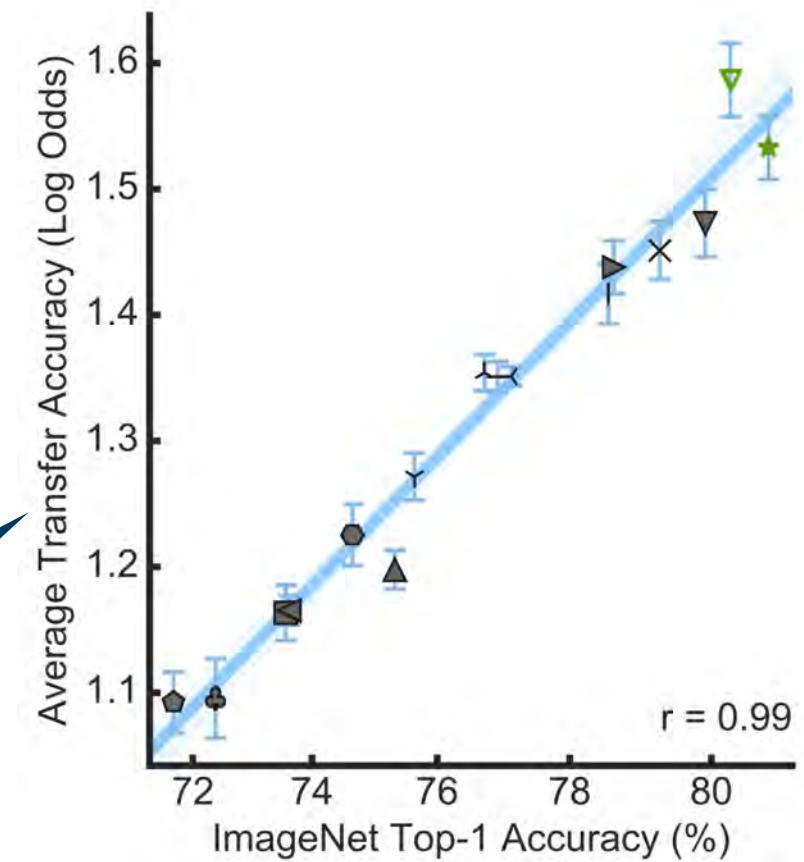
Better ImageNet lead to improvements on downstream tasks

Do Better ImageNet Models Transfer Better?

Simon Kornblith, Jacob Eisenstein, and Quoc V. Le
Google Brain
{skornblith, shlens, qvl}@google.com

Yes!

Tested on 12 different image classification sets



Computer vision tasks

**Image
Classification**



Cat

No spatial extent

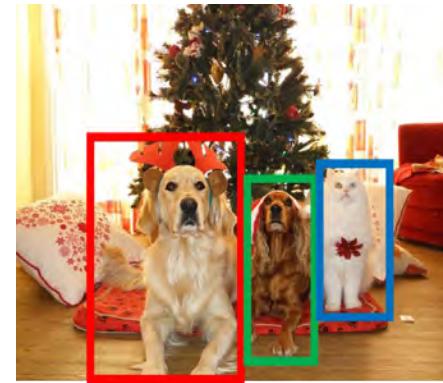
**Semantic
Segmentation**



Sky, Trees, Grass, Cat

No objects, just pixels

**Object
Detection**



Dog, Dog, Cat

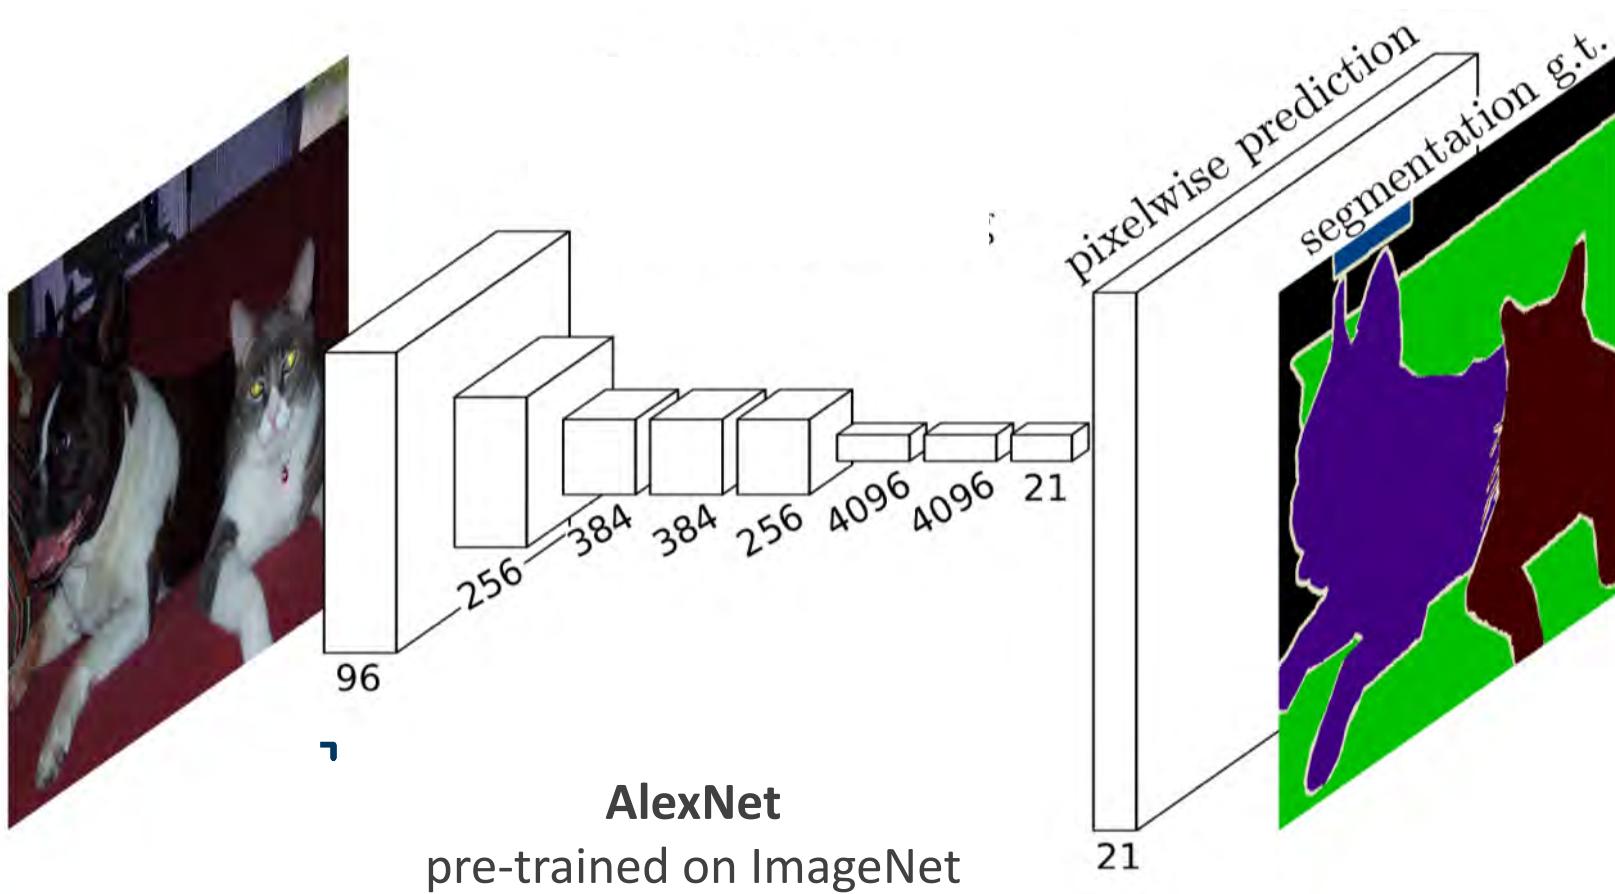
Multiple objects

**Instance
Segmentation**

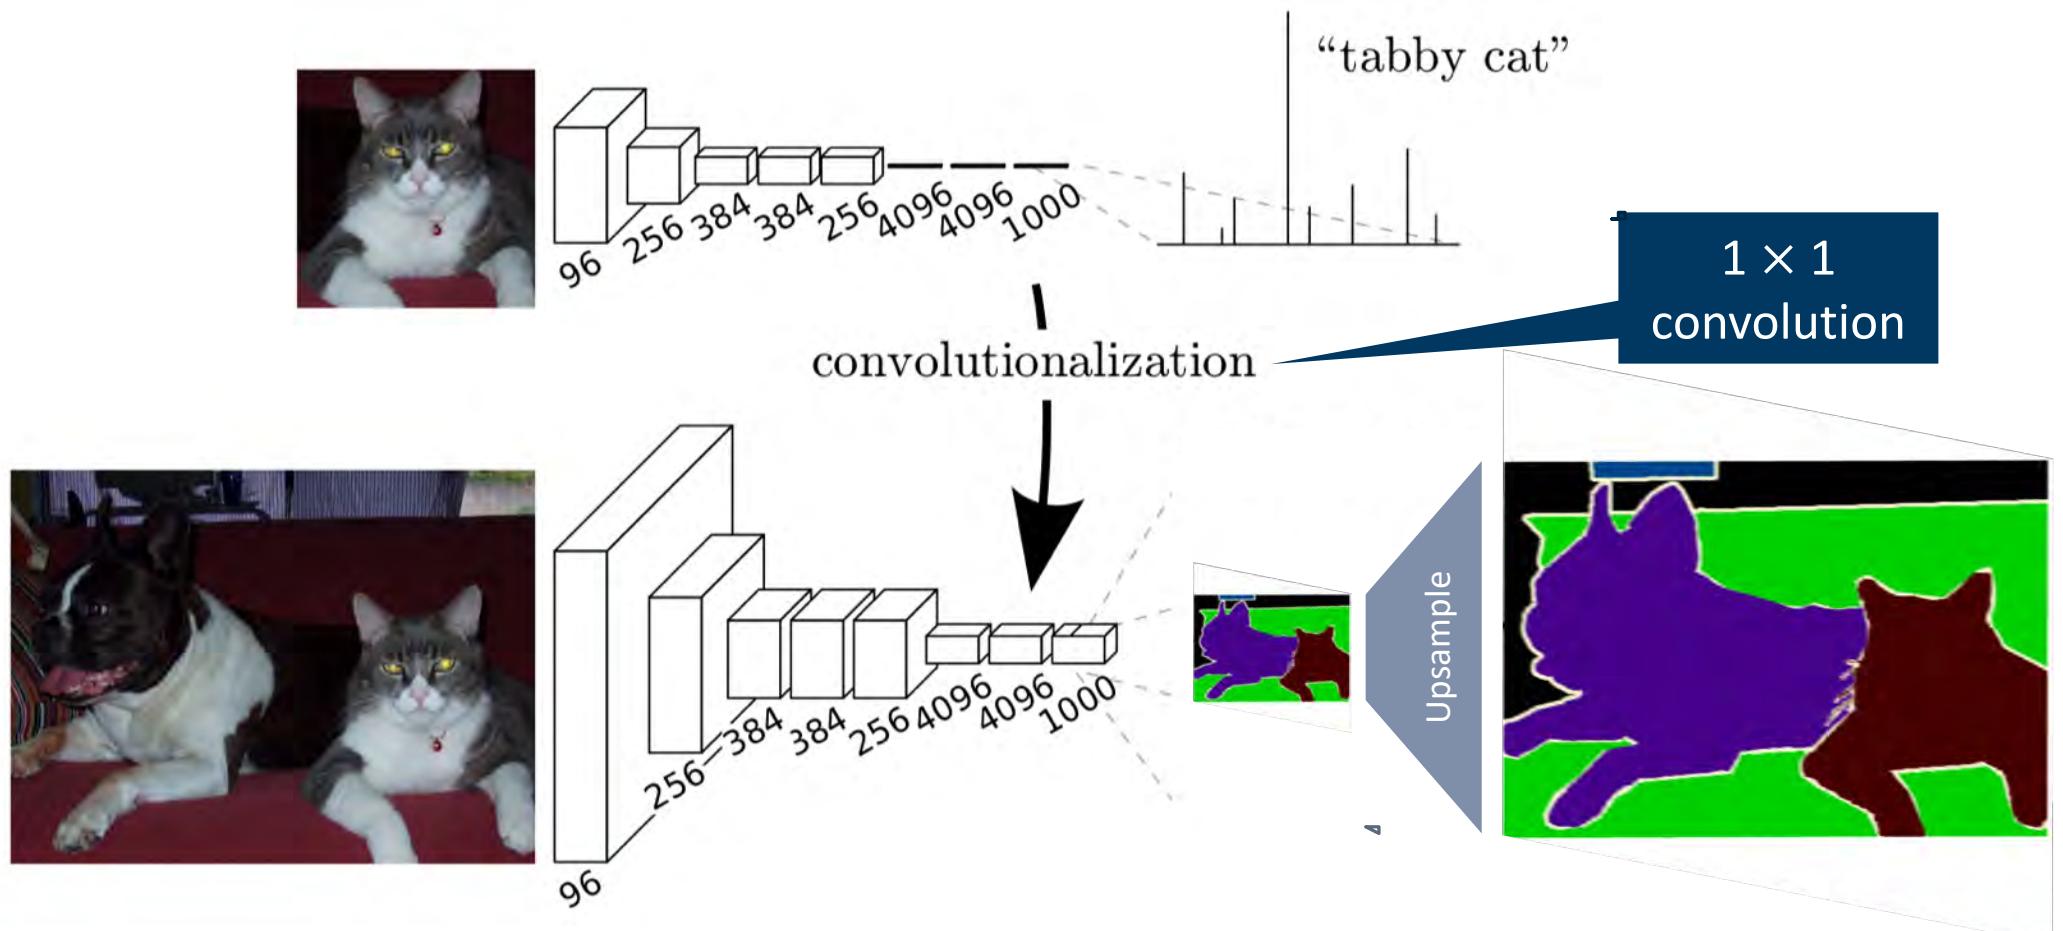


Dog, Dog, Cat

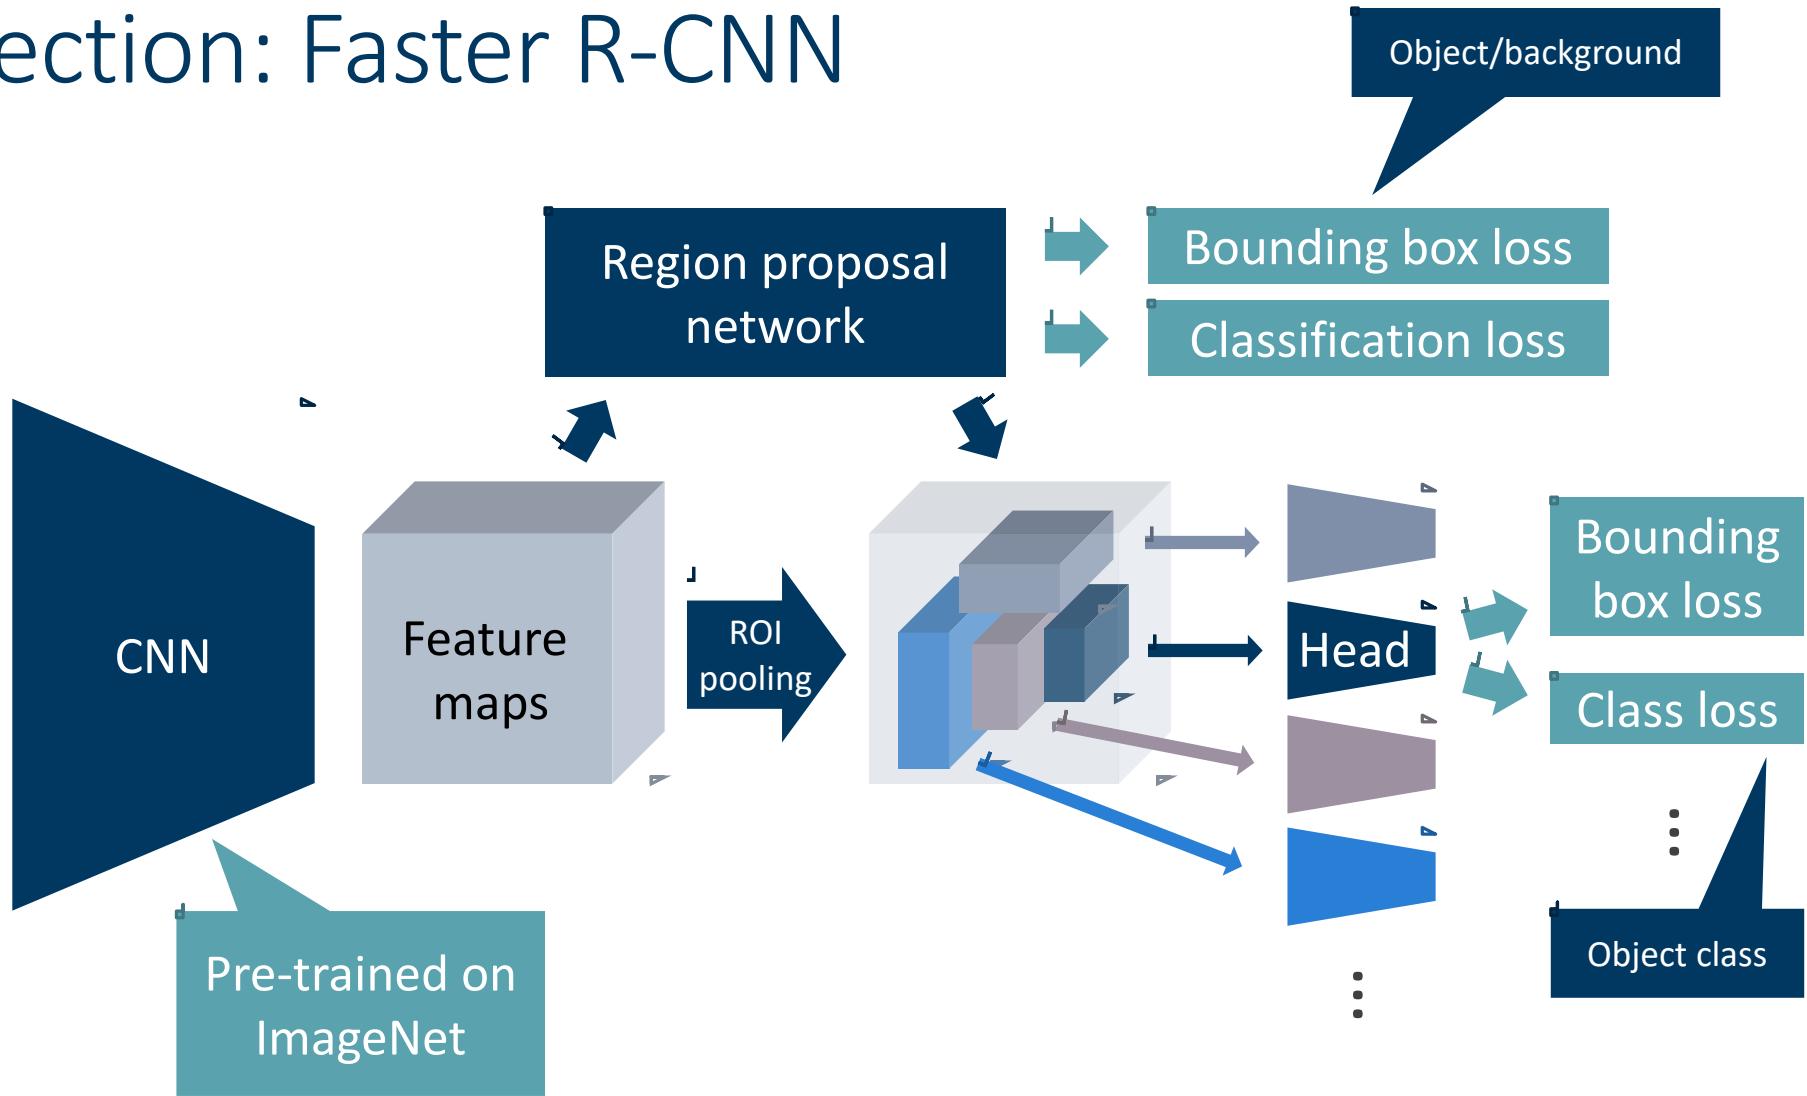
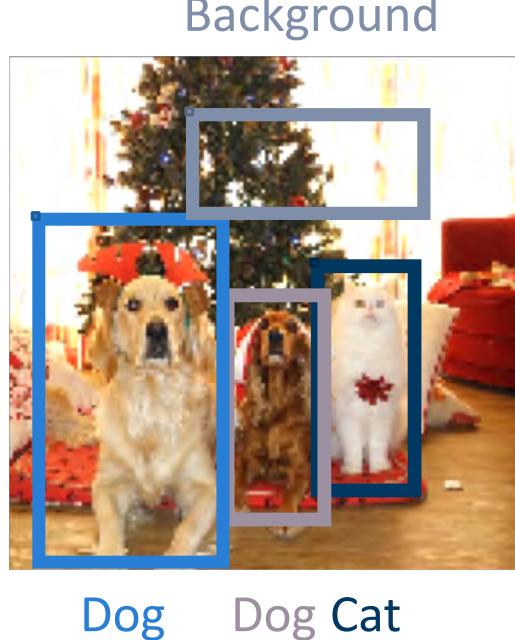
Semantic segmentation



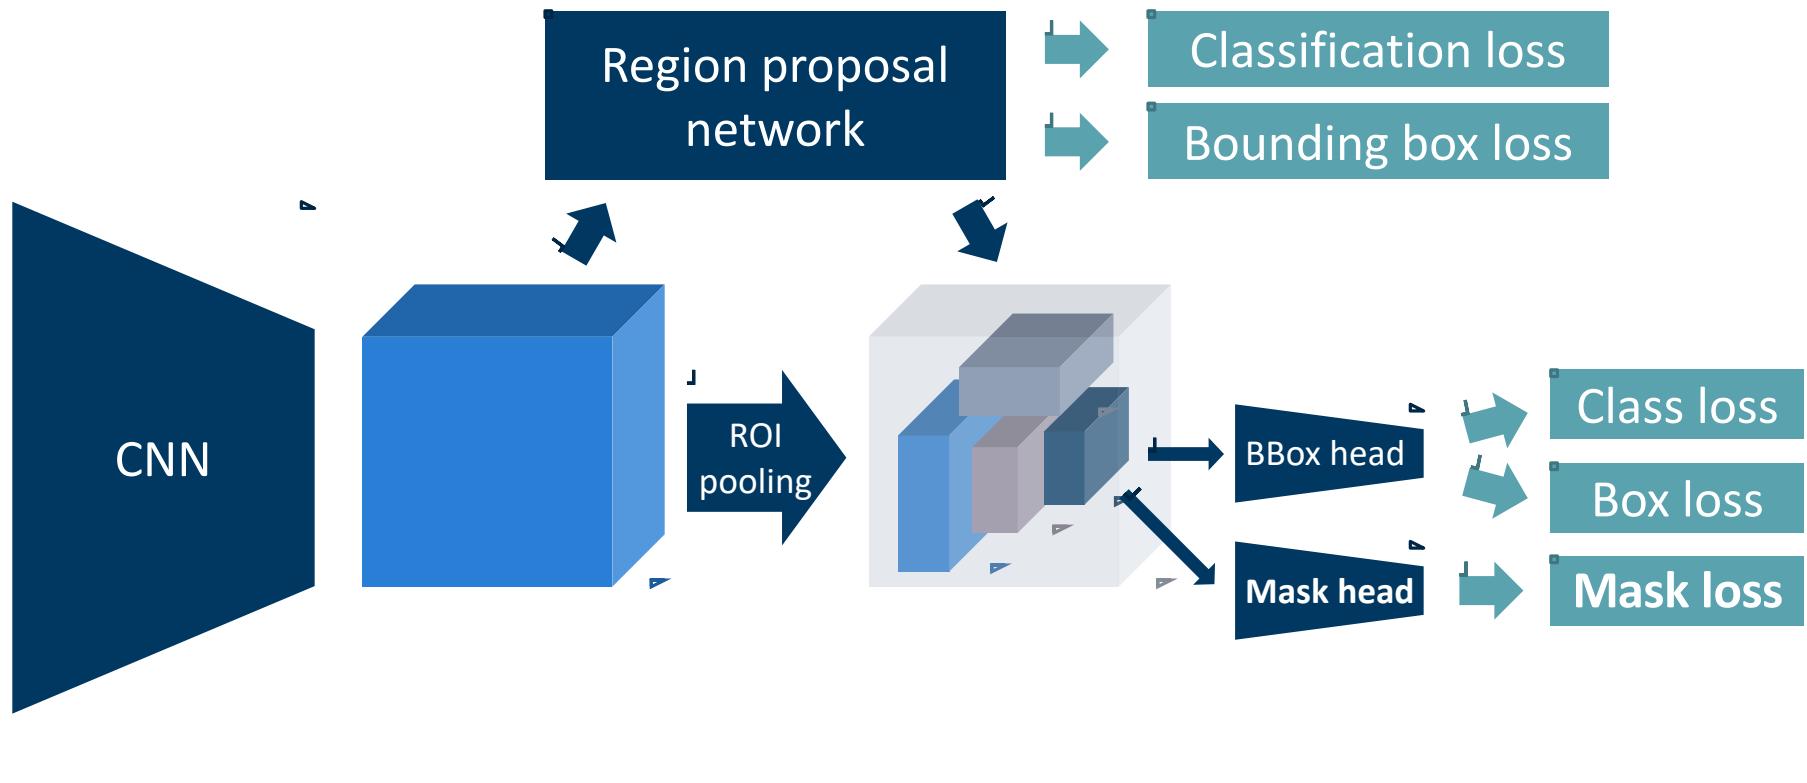
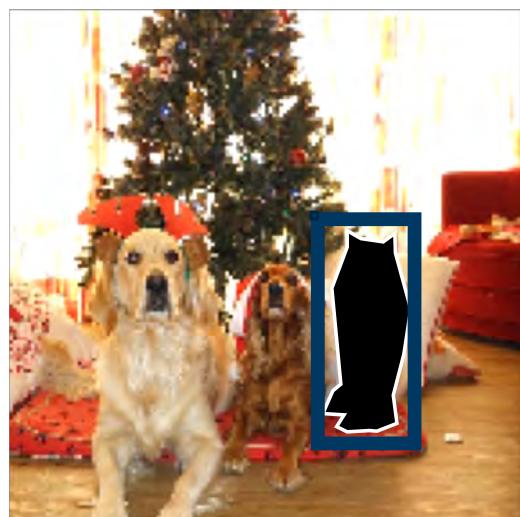
Turning fully-connected layers into conv layers



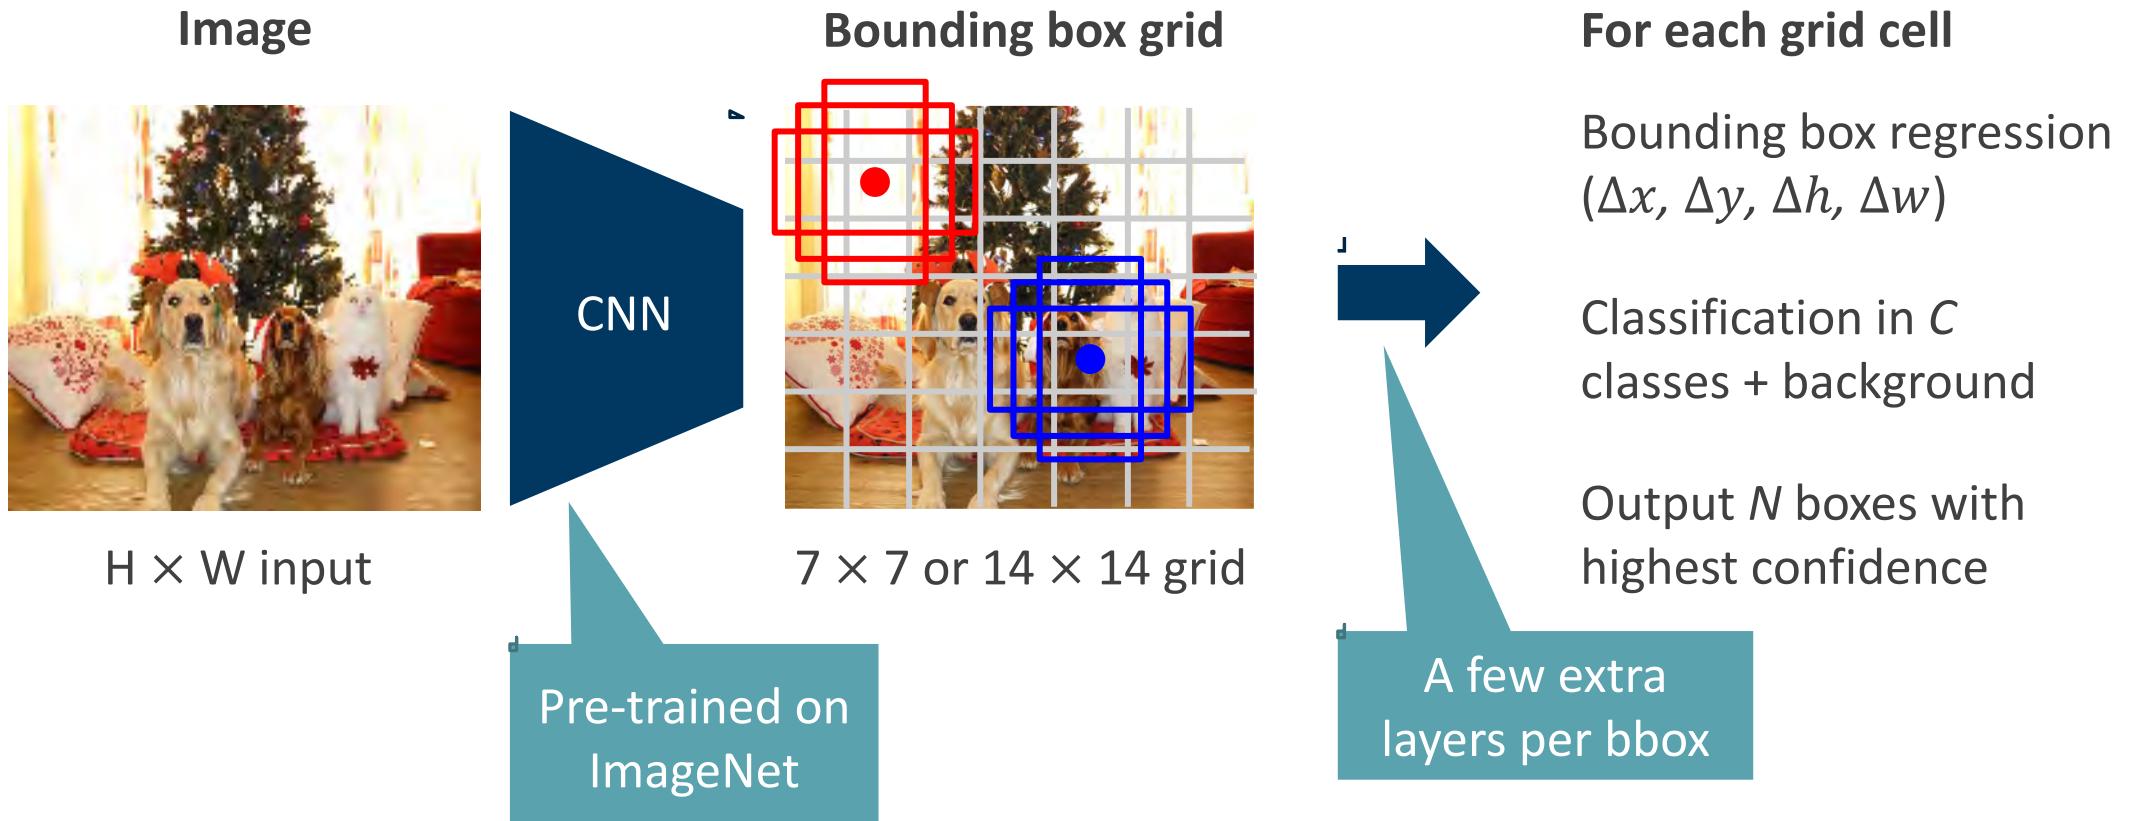
Object detection: Faster R-CNN



Instance segmentation: Mask R-CNN



Single-stage detectors: SSD, YOLO, RetinaNet



Pre-training on ImageNet is everywhere ...

Human pose estimation

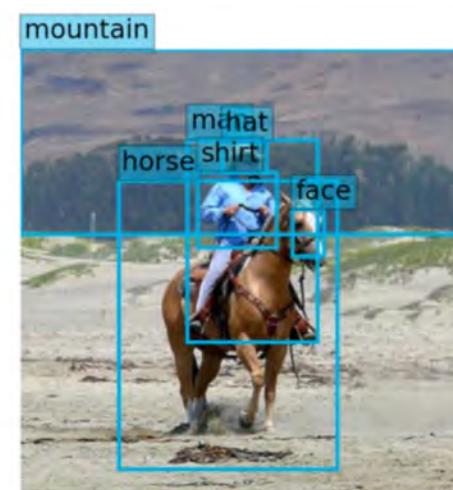


Image captioning



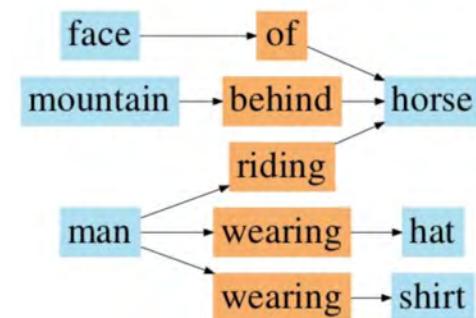
"two young girls are playing with
lego toy."

Scene graph prediction



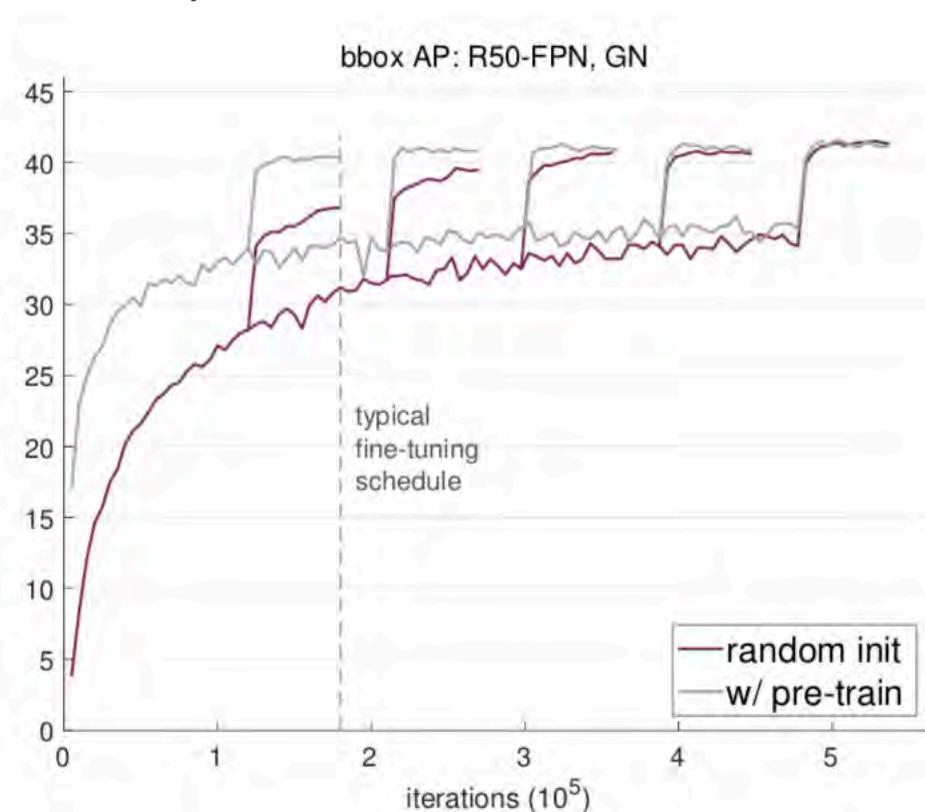
...

Many,
many
more



Transfer learning from ImageNet is the standard...

... but may not always be necessary



Object detection
on MS COCO dataset
123,287 images
886,284 instances

Visual saliency

PREDICTING WHERE PEOPLE LOOK

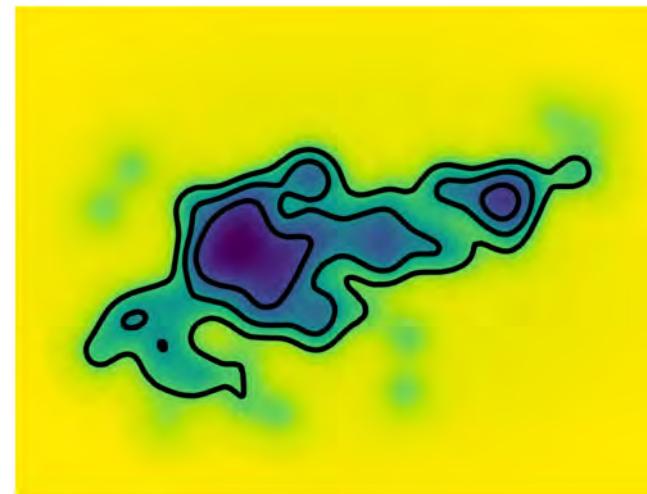
Image saliency: predicting where people look

Image shown



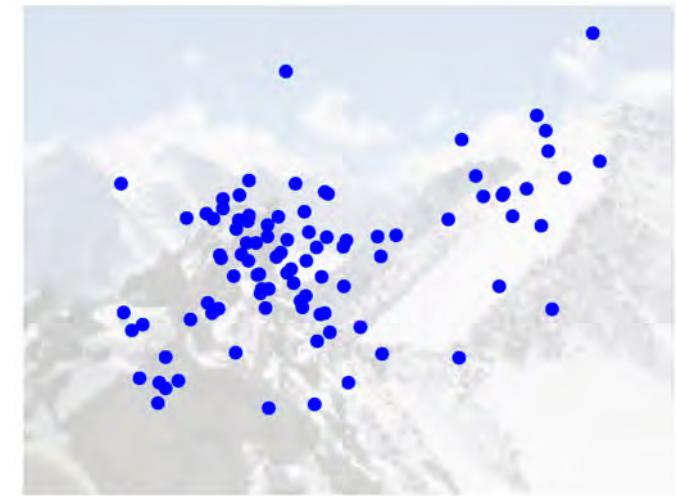
Predict

Goal:
Saliency map

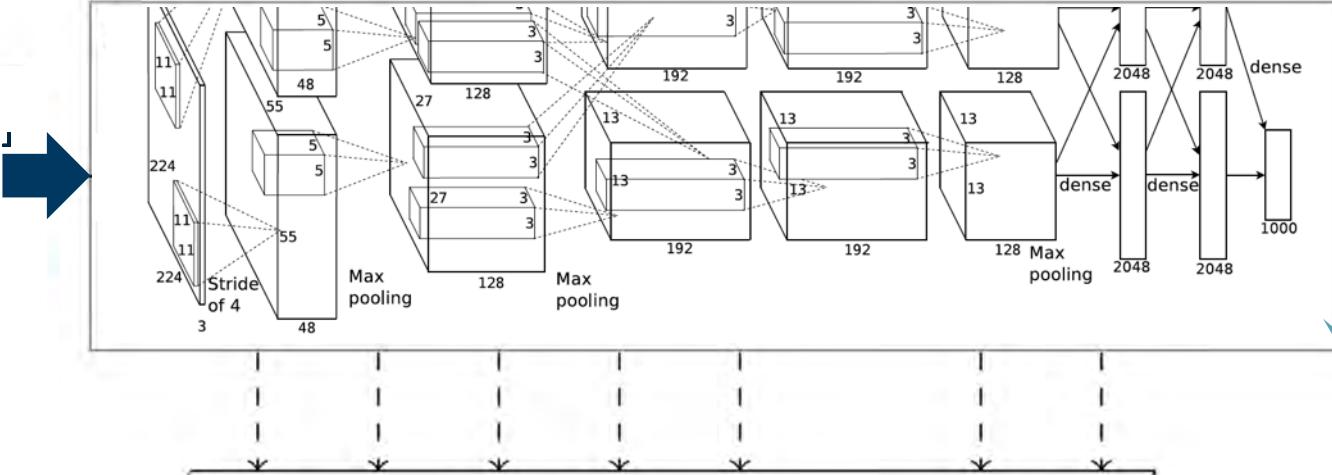


$$P(x, y \mid \text{image})$$

Measured eye fixations



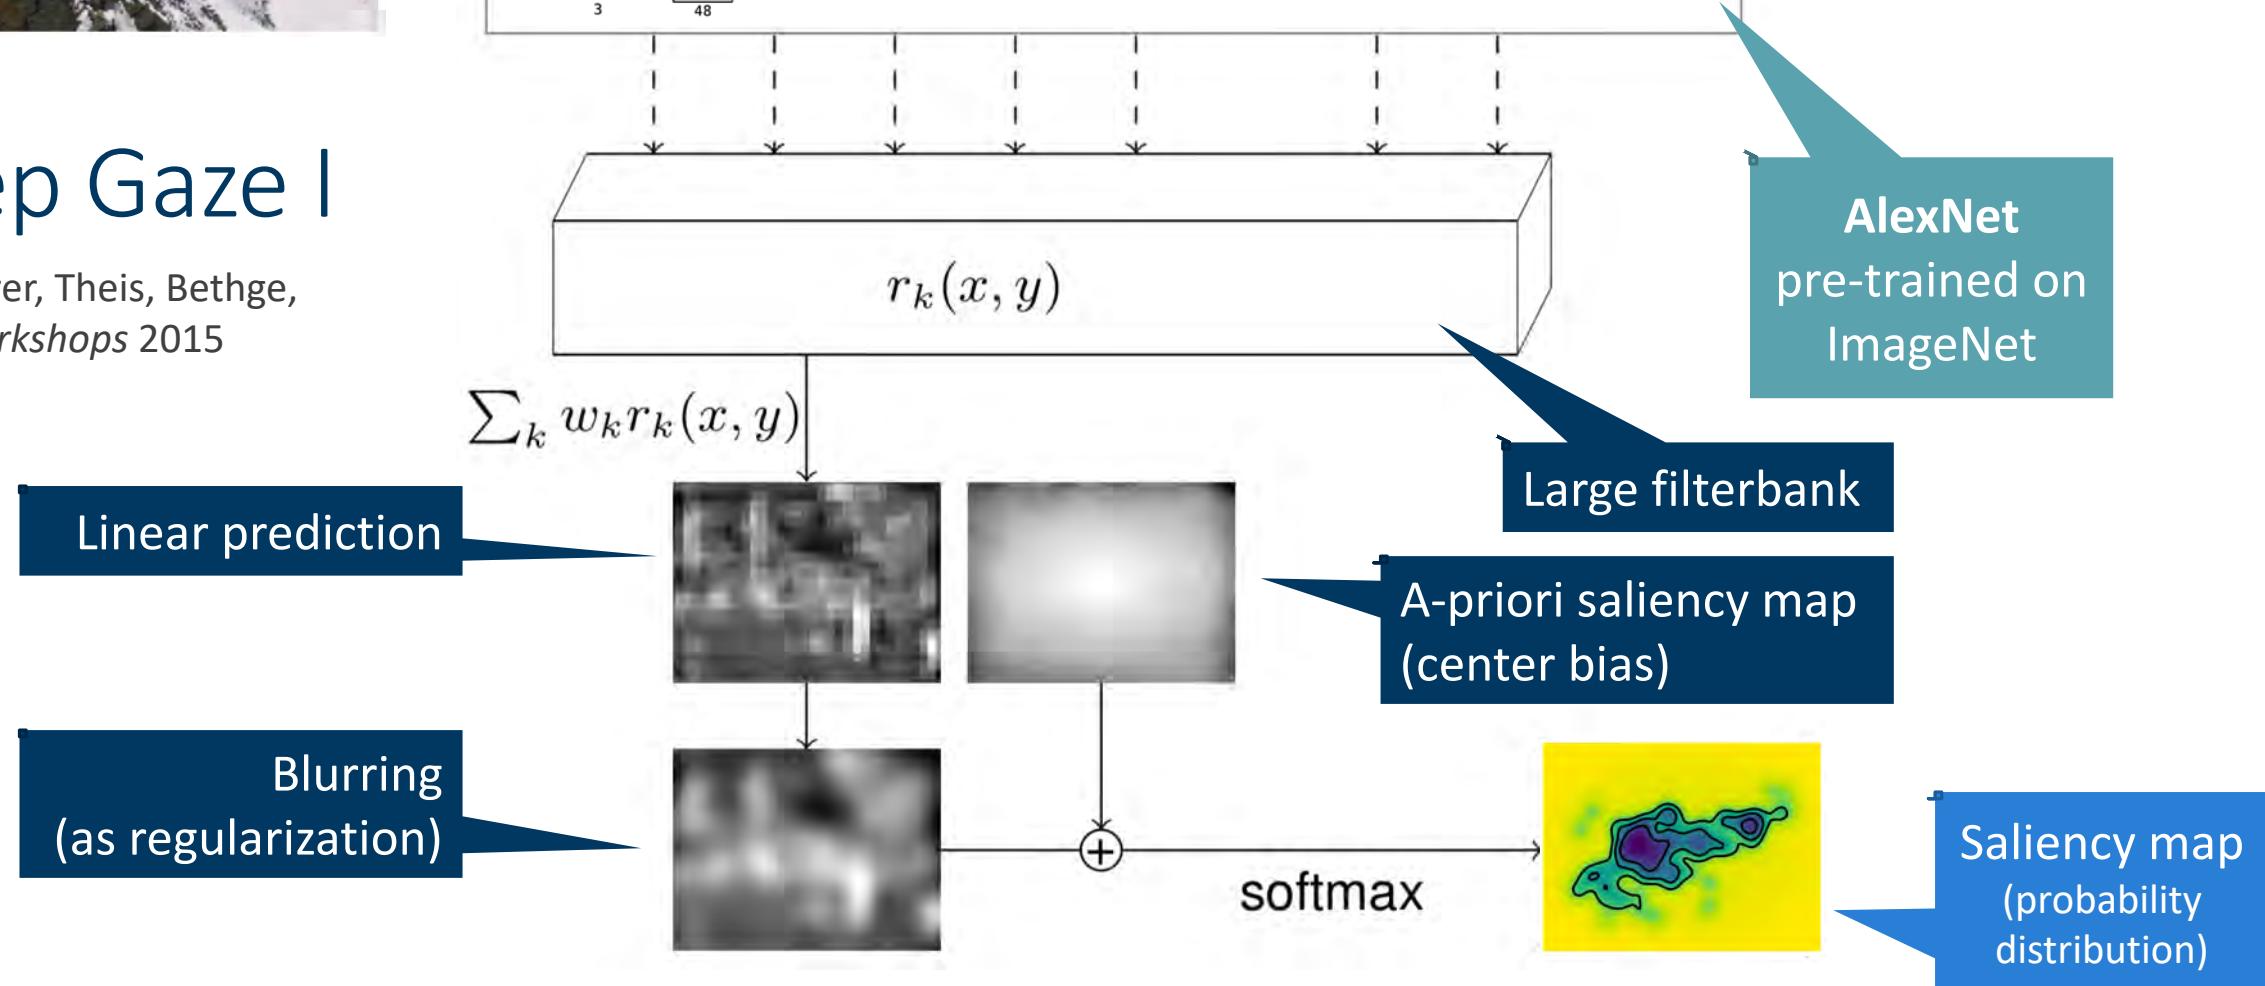
Loss



Deep Gaze I

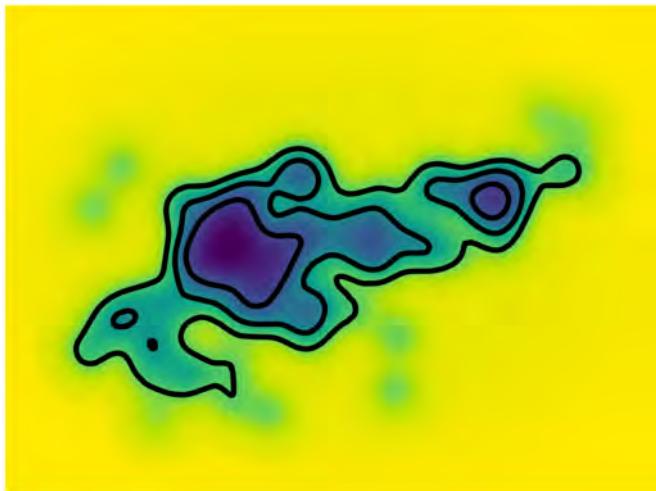
Kümmerer, Theis, Bethge,
ICLR Workshops 2015

AlexNet
pre-trained on
ImageNet



Deep Gaze: loss function

Saliency map



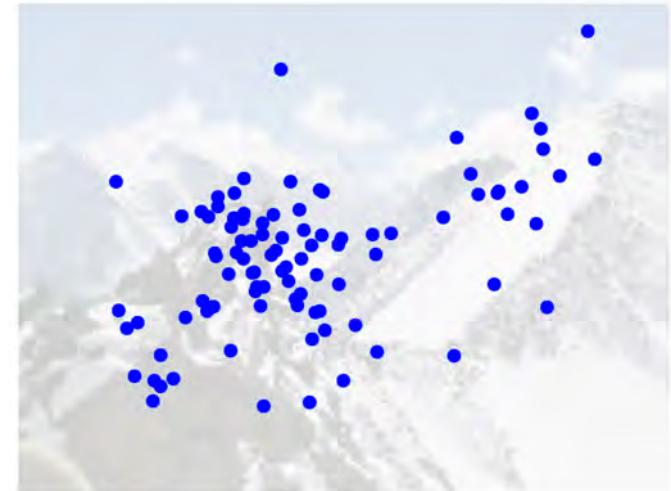
$P(x, y \mid \text{image})$

Loss:

negative log likelihood

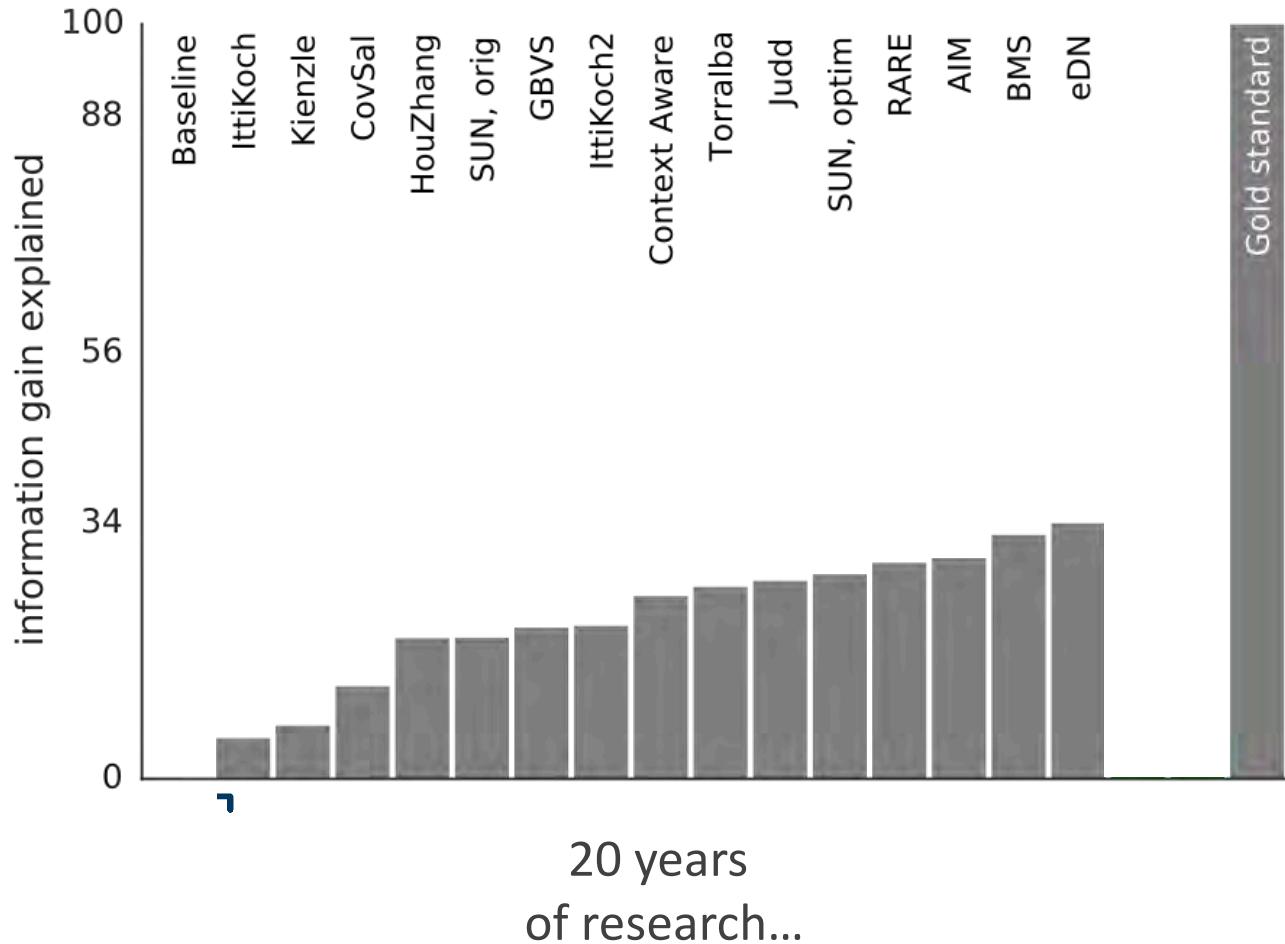
$$L = - \sum_{i=1}^N \log P(x_i, y_i \mid \text{image})$$

Measured eye fixations

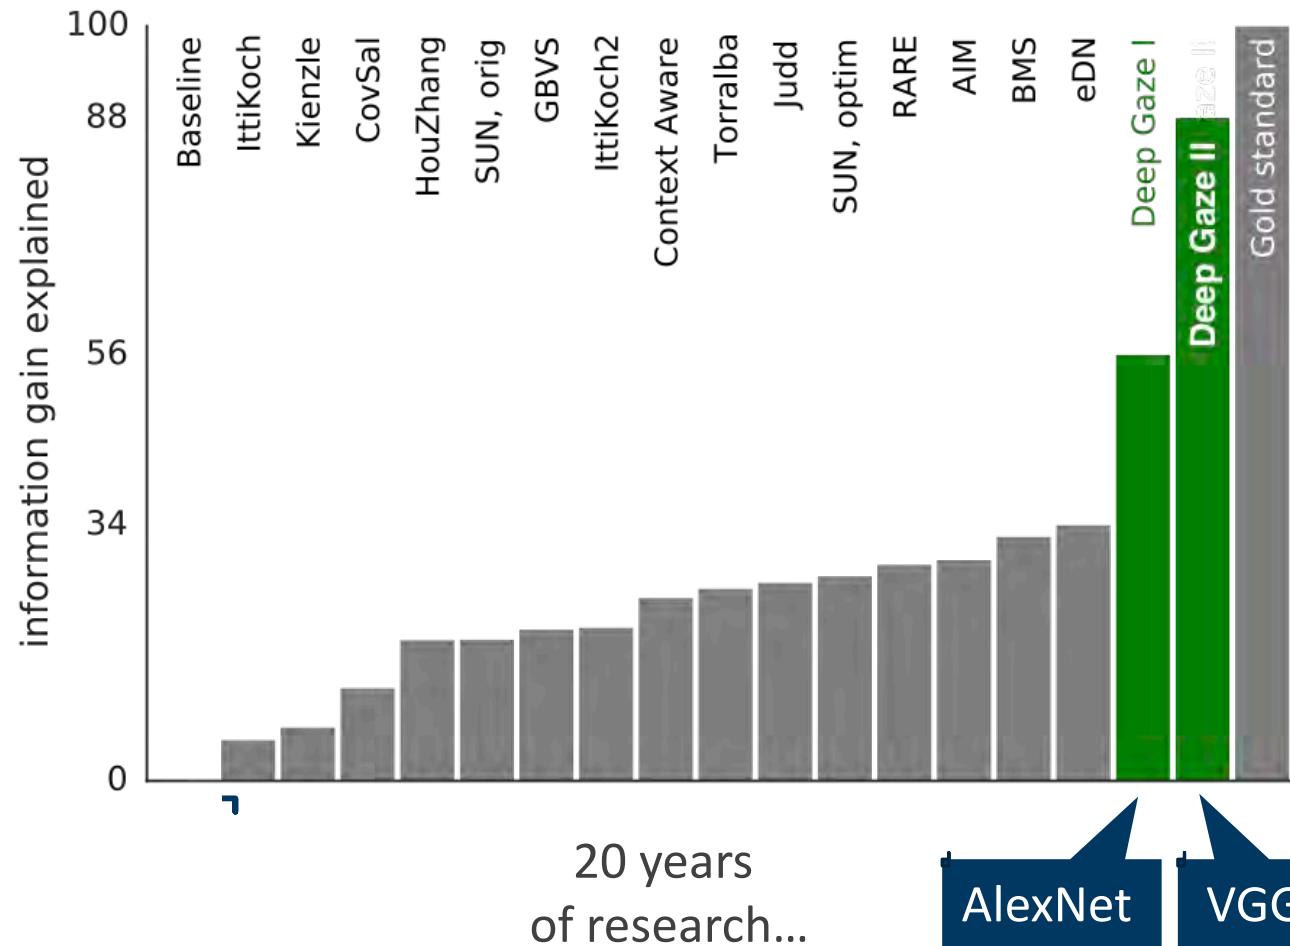


$\{(x_i, y_i)\}_{i=1}^N$

Deep Gaze I + II



Deep Gaze I + II



Kümmerer, Theis, Bethge,
ICLR Workshops 2015

Kümmerer, Wallis, Bethge, *arXiv 2016*

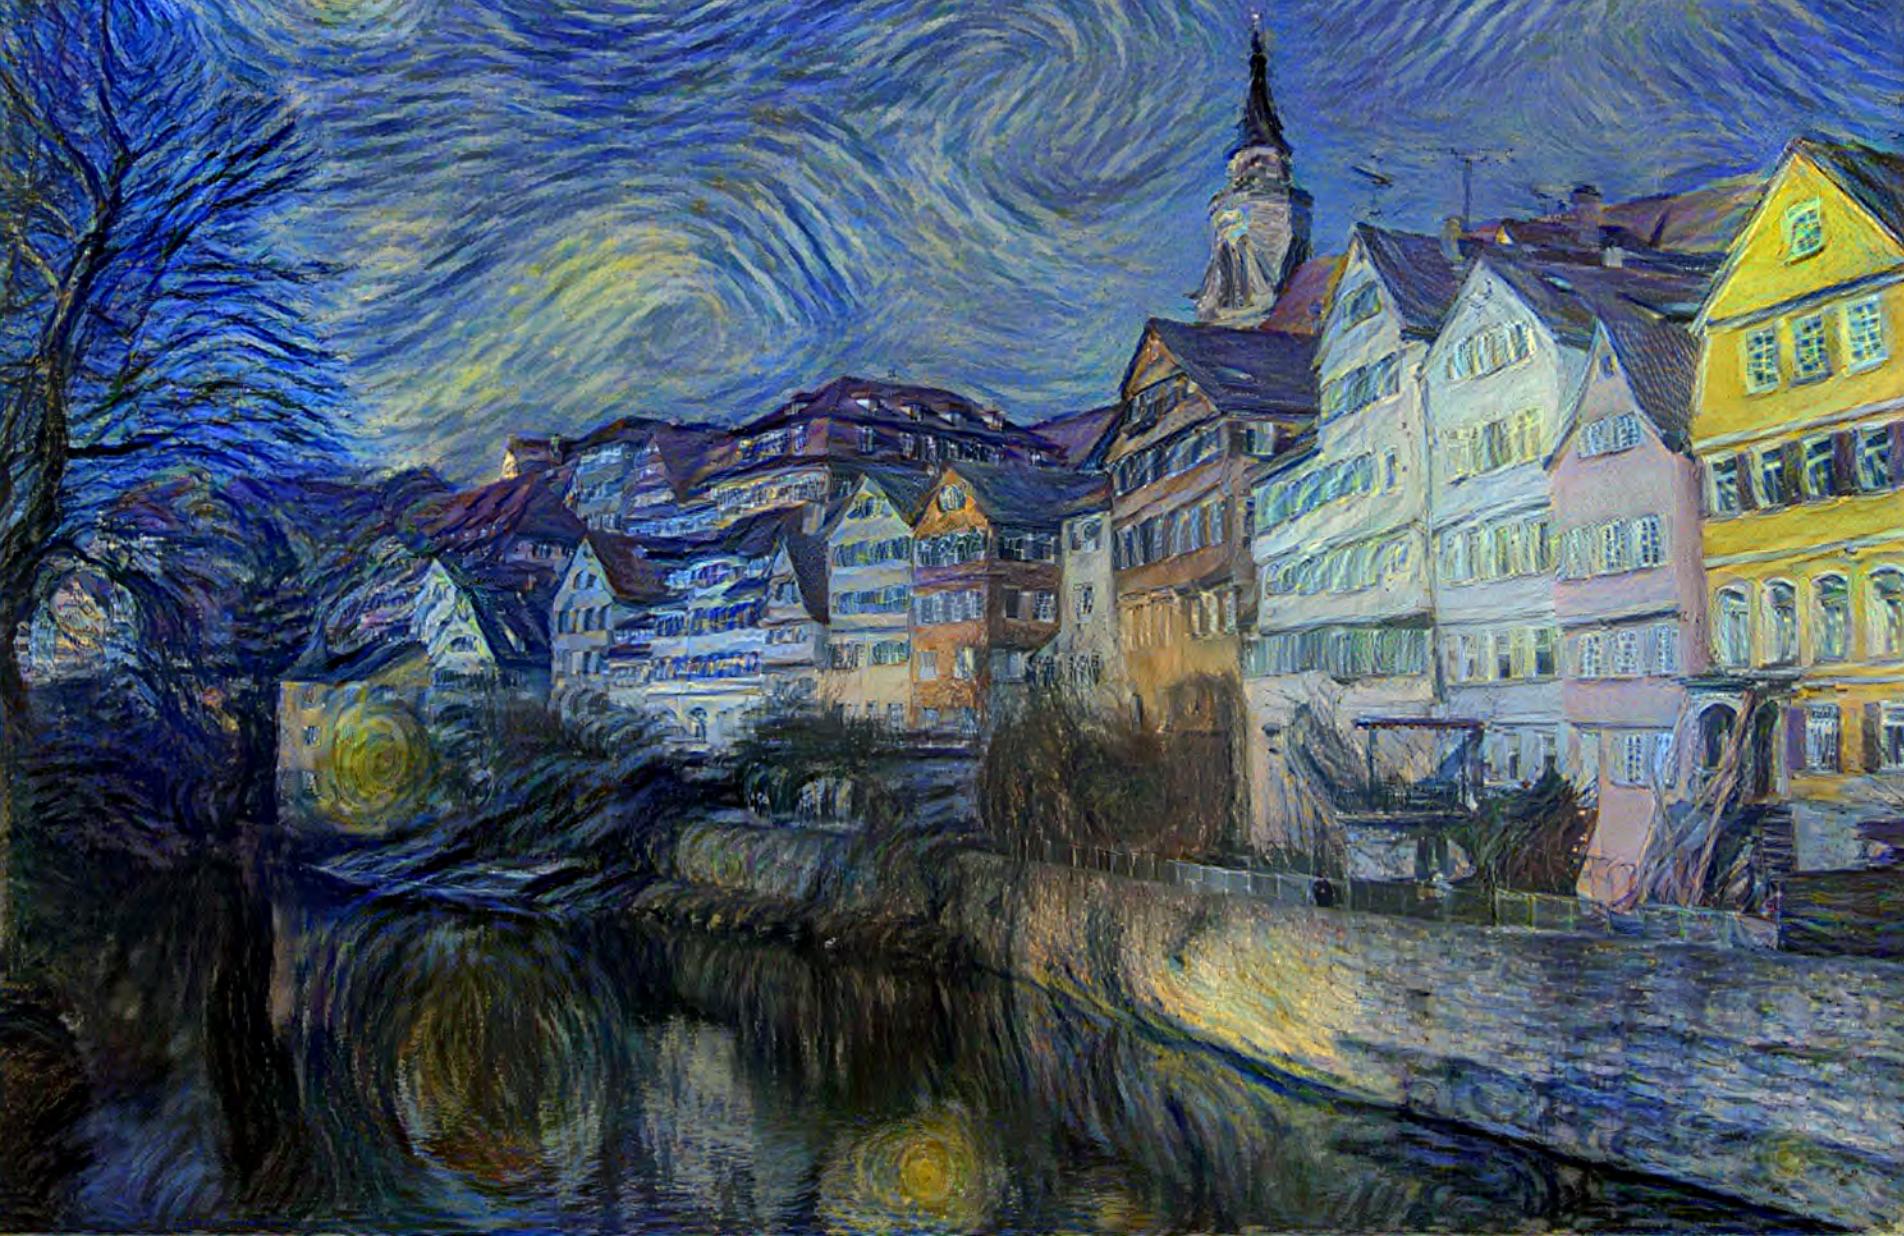
Neural style transfer

TEXTURE SYNTHESIS AND PERCEPTUAL LOSS FUNCTIONS

History of art – Tübingen



Vincent van Gogh (1889)



William Turner (1805)



Pablo Picasso (1910)



Edvard Munch (1893)



Wassily Kandinsky (1913)



A Neural Algorithm of Artistic Style



Gatys, Ecker, Bethge (CVPR 2016)



Neural Style Transfer

Content



Style (→ texture)



A brief history of parametric texture modeling

Julesz (1962)

- Conjecture: N^{th} order summary statistics determine texture perception

Portilla & Simoncelli (2000)

- Summary statistics of features from early visual system (steerable pyramid)

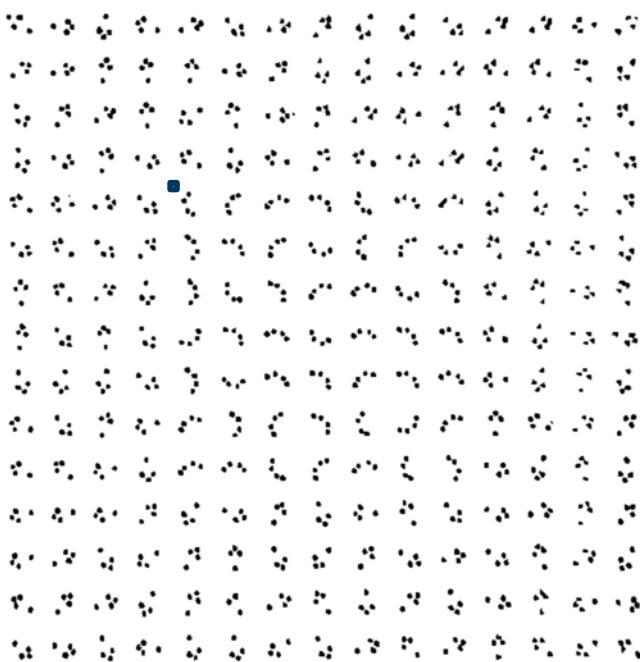
Gatys, Ecker Bethge (NIPS 2015)

- Summary statistics of features from CNNs trained on ImageNet

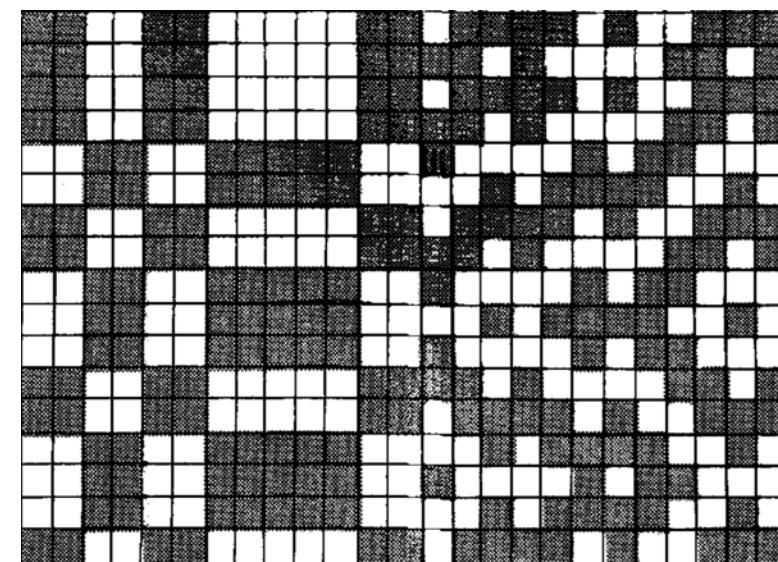
Julesz' conjecture (1962)

N^{th} order summary statistics determine texture perception.

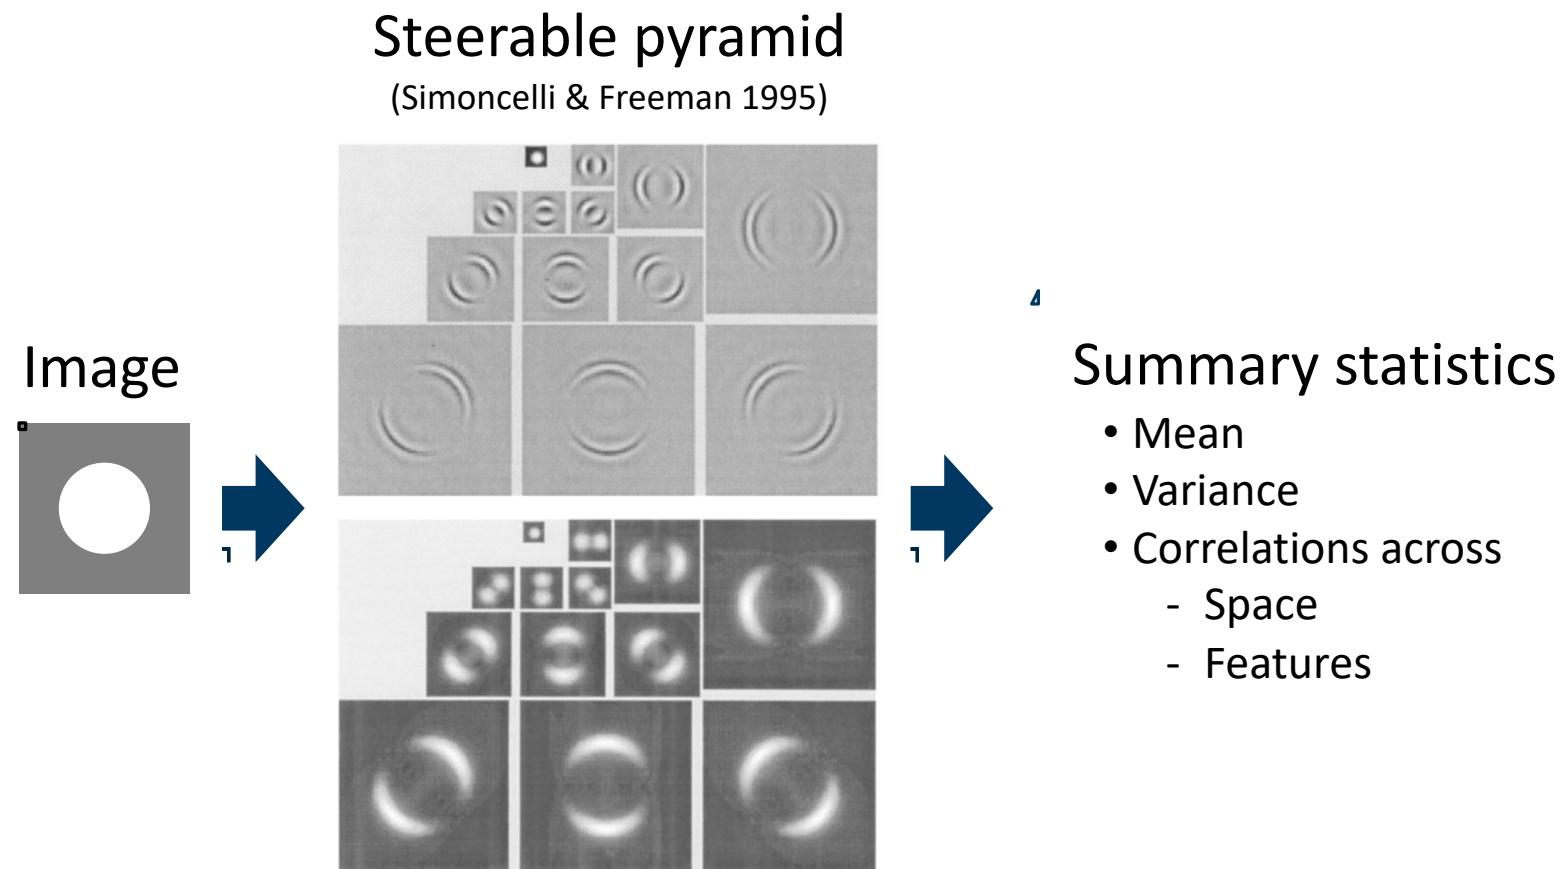
Uniform 2nd order statistics
(power spectrum)



Uniform 3rd order statistics

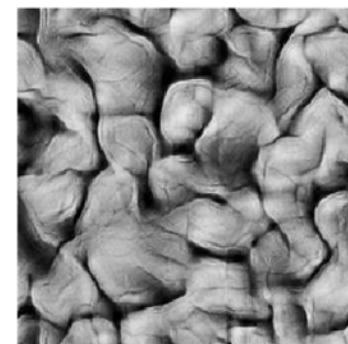
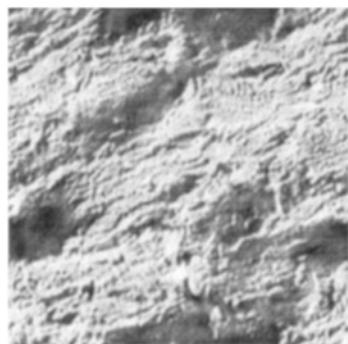
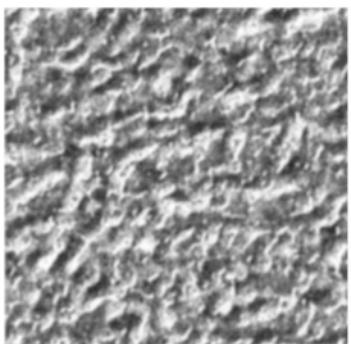
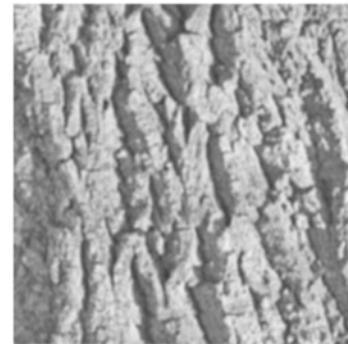
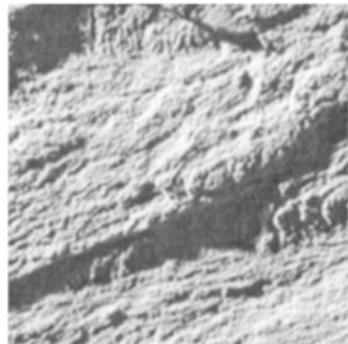
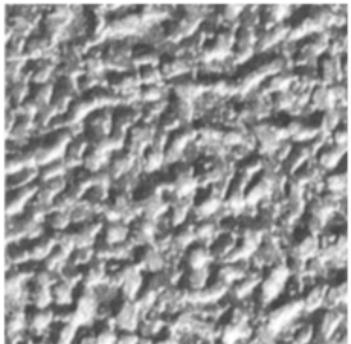


Portilla & Simoncelli texture model



Portilla & Simoncelli: results

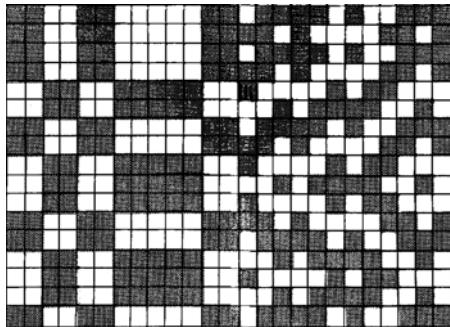
Original images



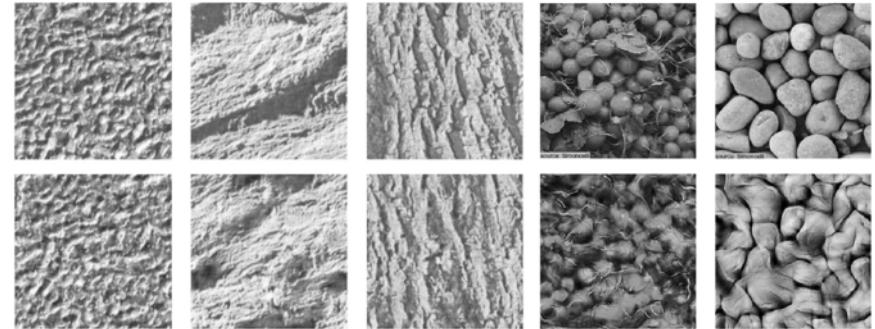
Synthesized by Portilla & Simoncelli

Portilla & Simoncelli (2000)

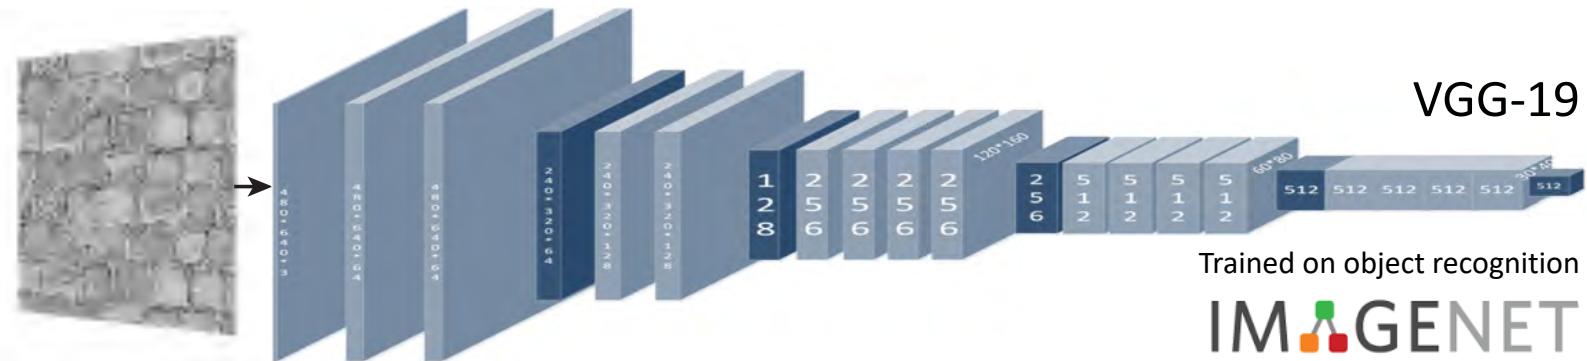
Texture models: summary statistics



Julesz (1962):
 N^{th} order pixel stats

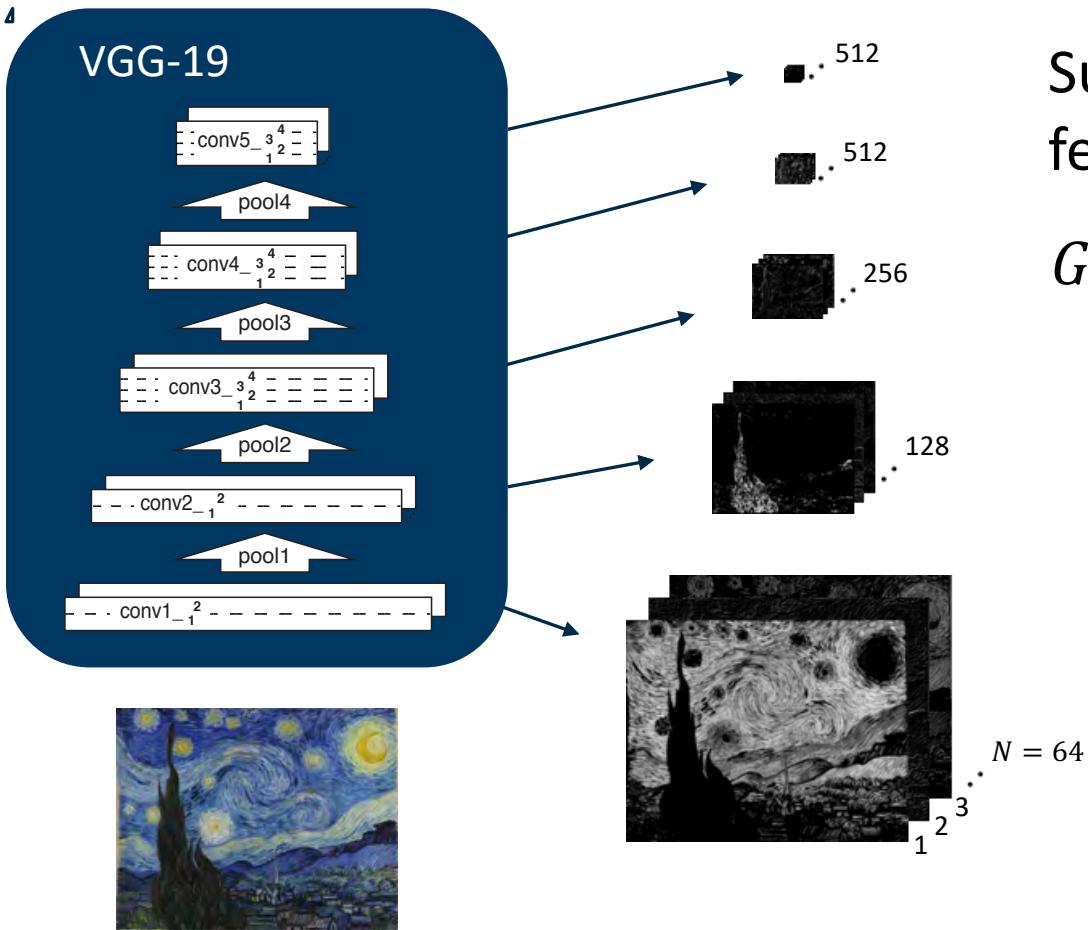


Portilla & Simoncelli (2000):
V1 summary statistics → Steerable Pyramid



Gatys, Ecker, Bethge (2015): Entire visual pathway → CNN features

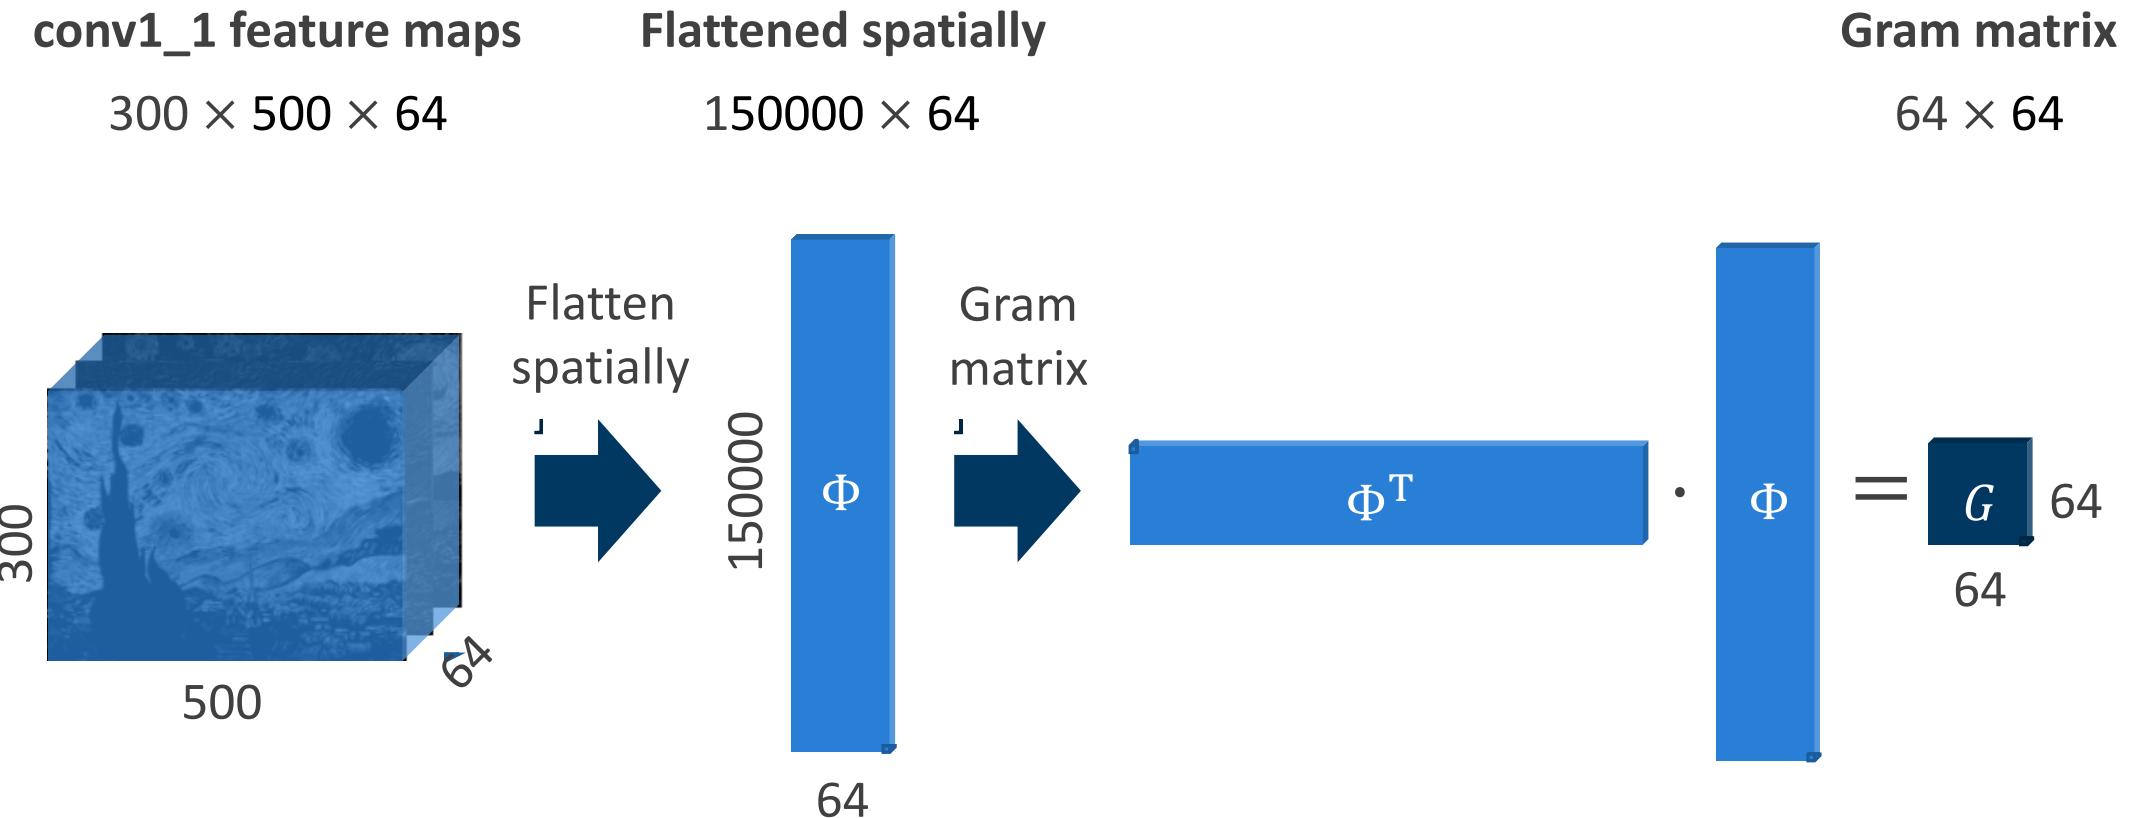
CNN: Multiscale nonlinear filter bank



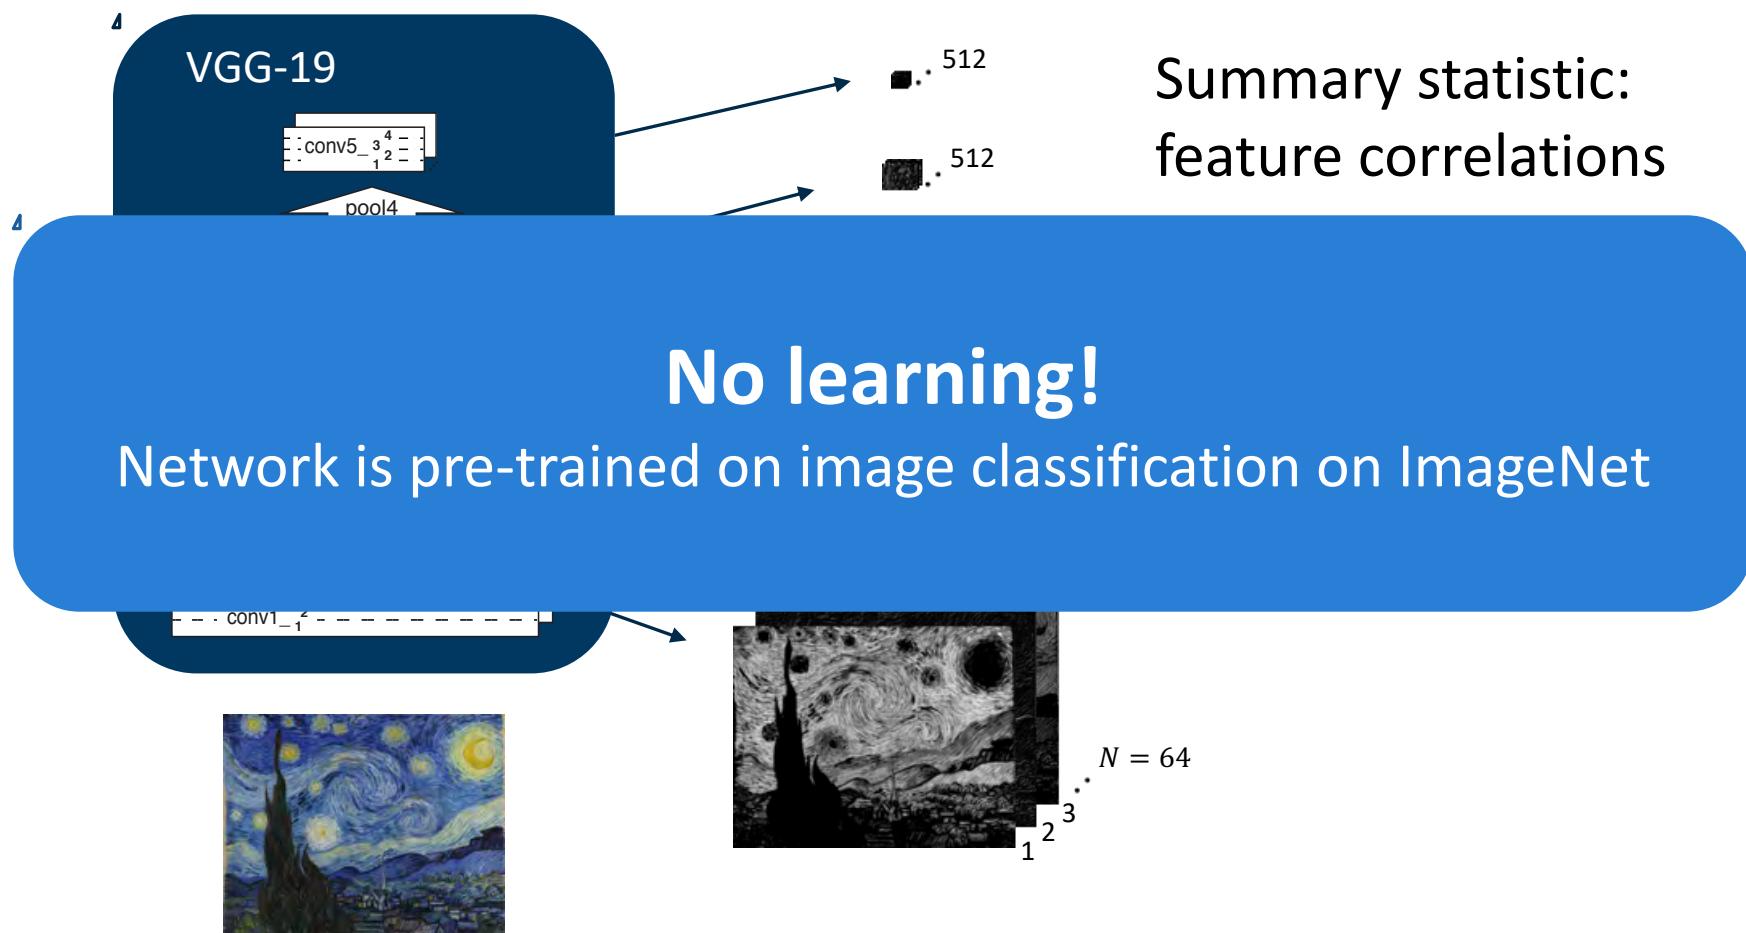
Summary statistic:
feature correlations

$$G = \Phi^T \Phi$$

Gram matrix computation

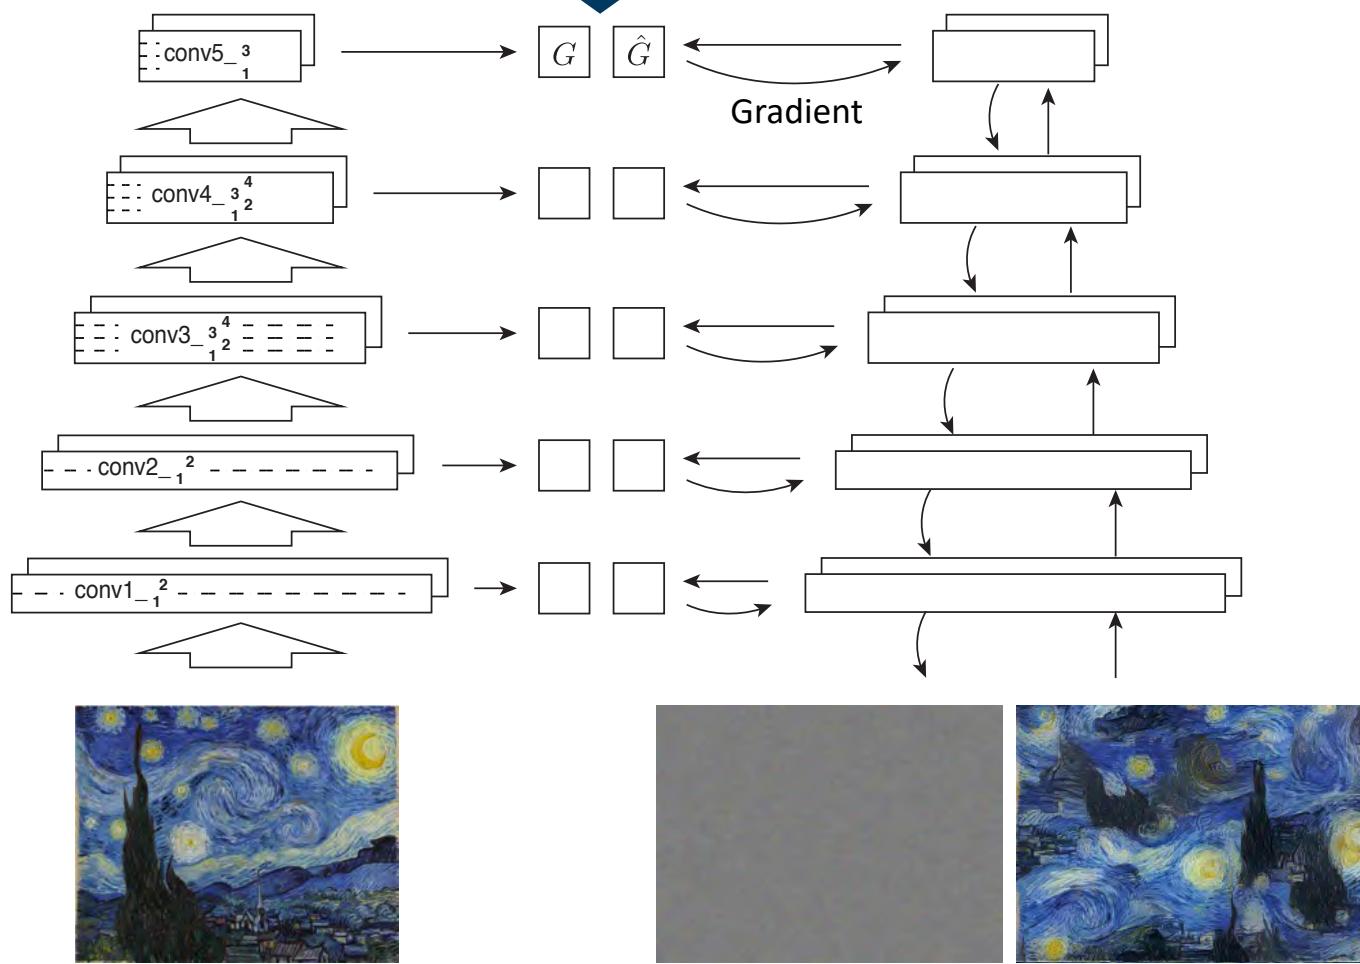


CNN: Multiscale nonlinear filter bank



Texture loss

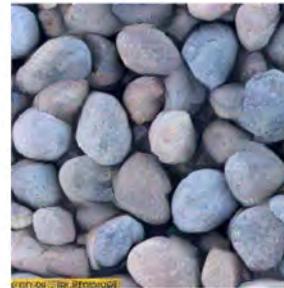
$$\mathcal{L}_{\text{style}} = \sum_{ij} (G_{ij} - \hat{G}_{ij})^2$$



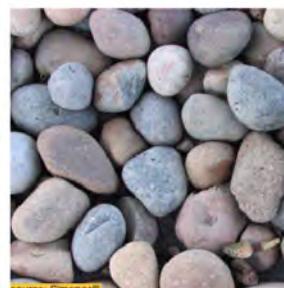
Gradient descent

Textures based on deep features

CNN-based model



Original image

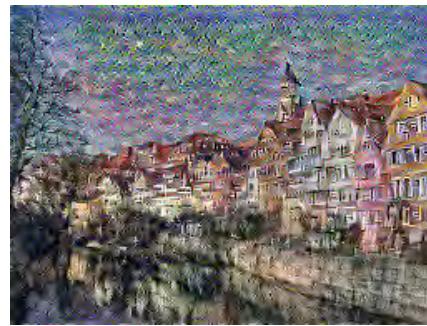


Portilla & Simoncelli



Neural Style Transfer

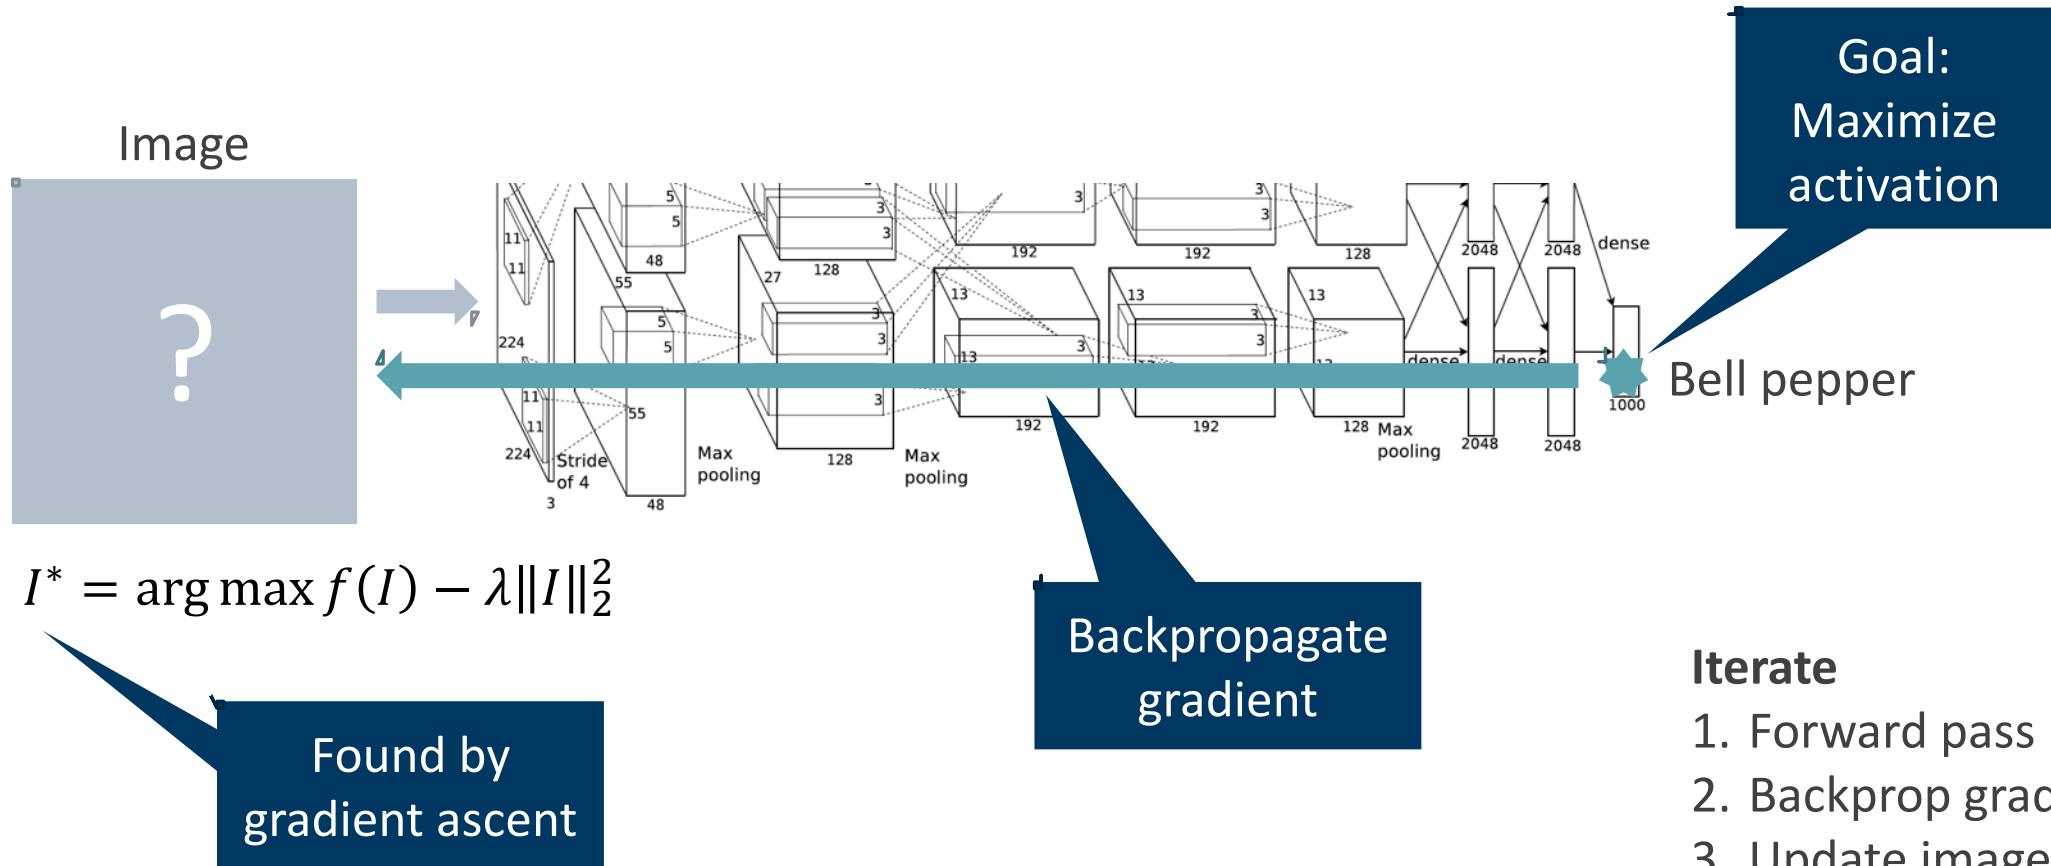
Content



Style (→ texture)



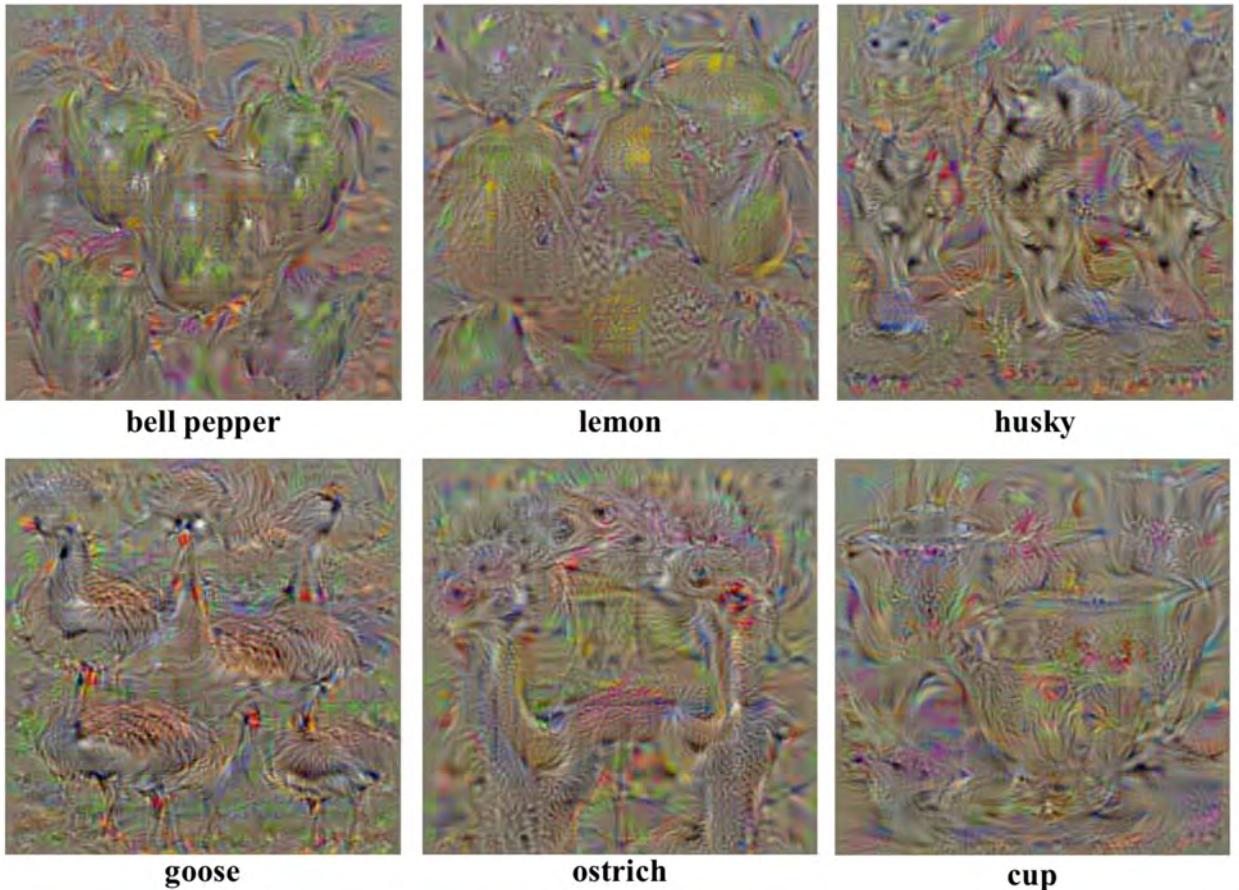
Visualizing CNN features by activity maximization



Visualizing classification layer

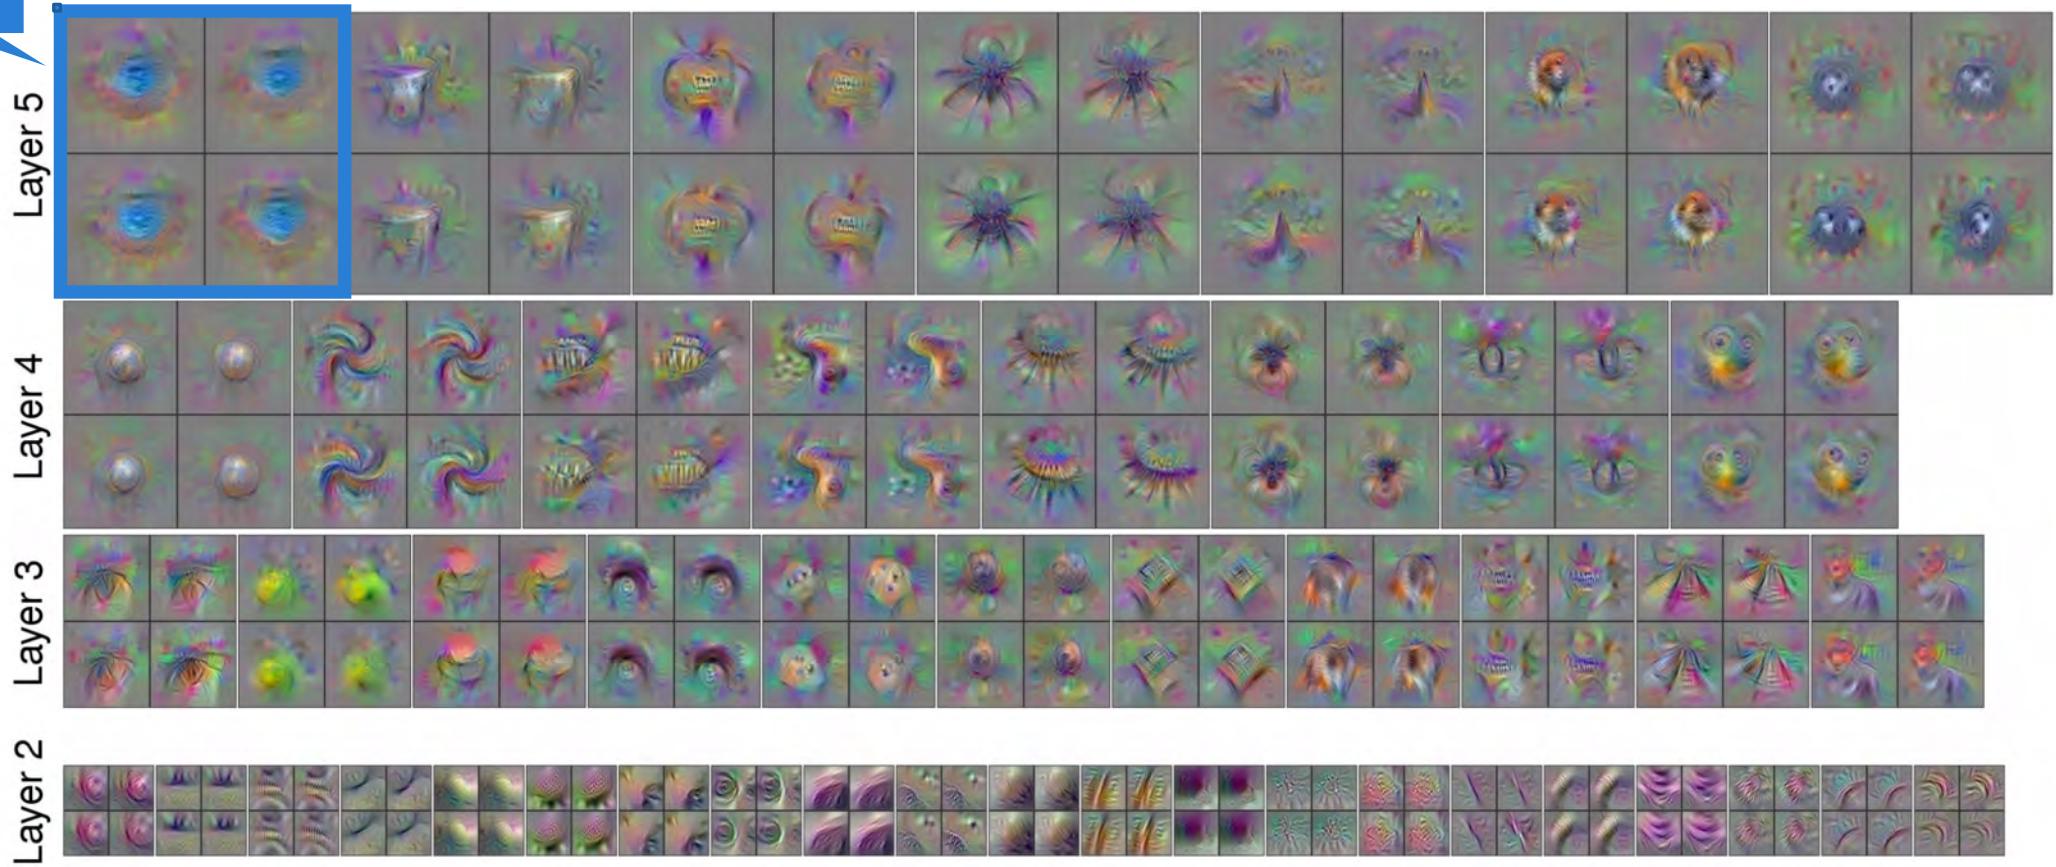
$$I^* = \arg \max f(I) - \lambda \|I\|_2^2$$

Regularizer to get
“natural” images



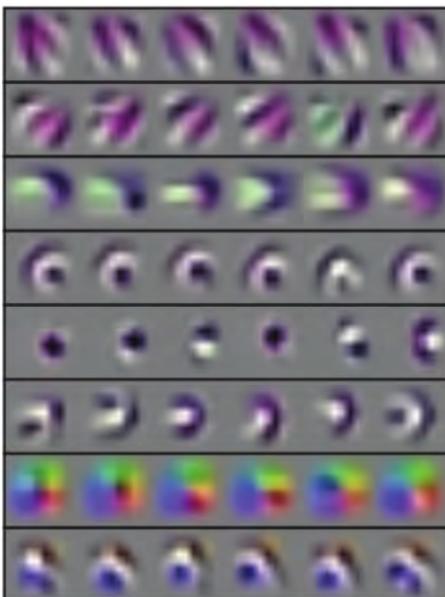
Visualization of intermediate layers of AlexNet

1 unit



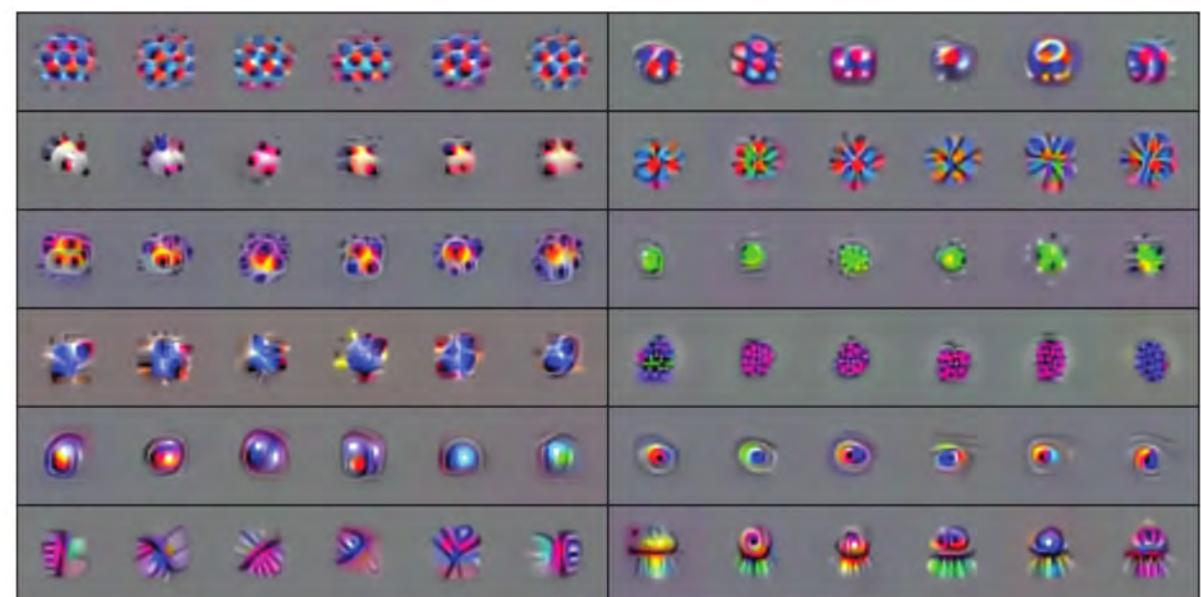
Visualization of intermediate layers of VGG-19

conv2_1



Early: selective

conv3_3



Intermediate/late: Increased invariance

Optimization target here: obtain diverse set of images that maximize overall activity

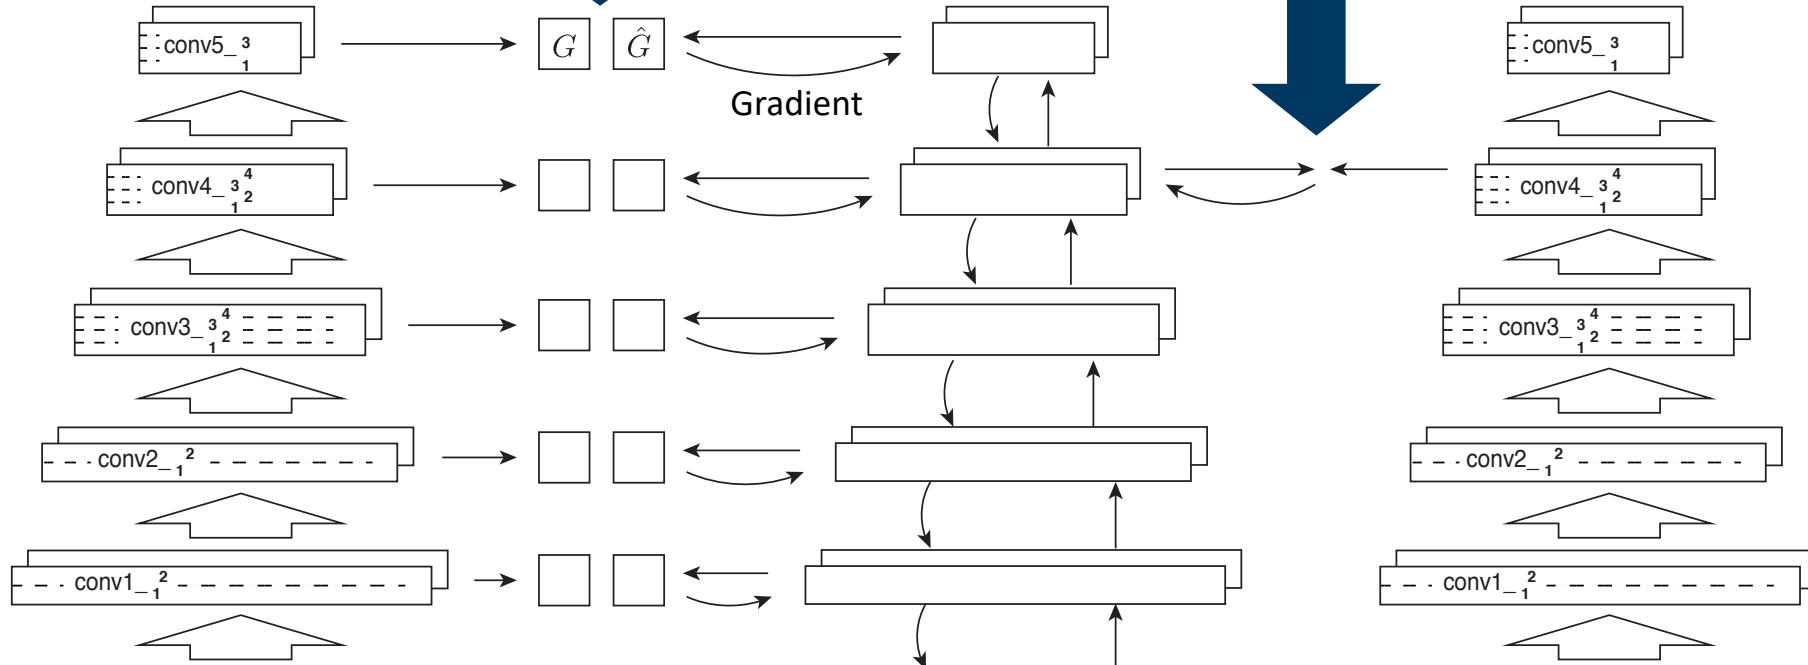
Style loss

$$\mathcal{L}_{\text{style}} = \sum_{ij} (G_{ij} - \hat{G}_{ij})^2$$

Content loss

$$\mathcal{L}_{\text{content}} = \sum_{ijk} (\Phi_{ijk} - \hat{\Phi}_{ijk})^2$$

Gradient



Gradient descent

Gatys, Ecker Bethge, CVPR 2016

Timelapse





Create your own artwork
<https://deepart.io>

TURN ANY PHOTO INTO AN ARTWORK – FOR FREE!

We use an algorithm inspired by the human brain. It uses the stylistic elements of one image to draw the content of another. Get your own artwork in just three steps.

1 Upload photo

The first picture defines the scene you would like to have painted.



2 Choose style

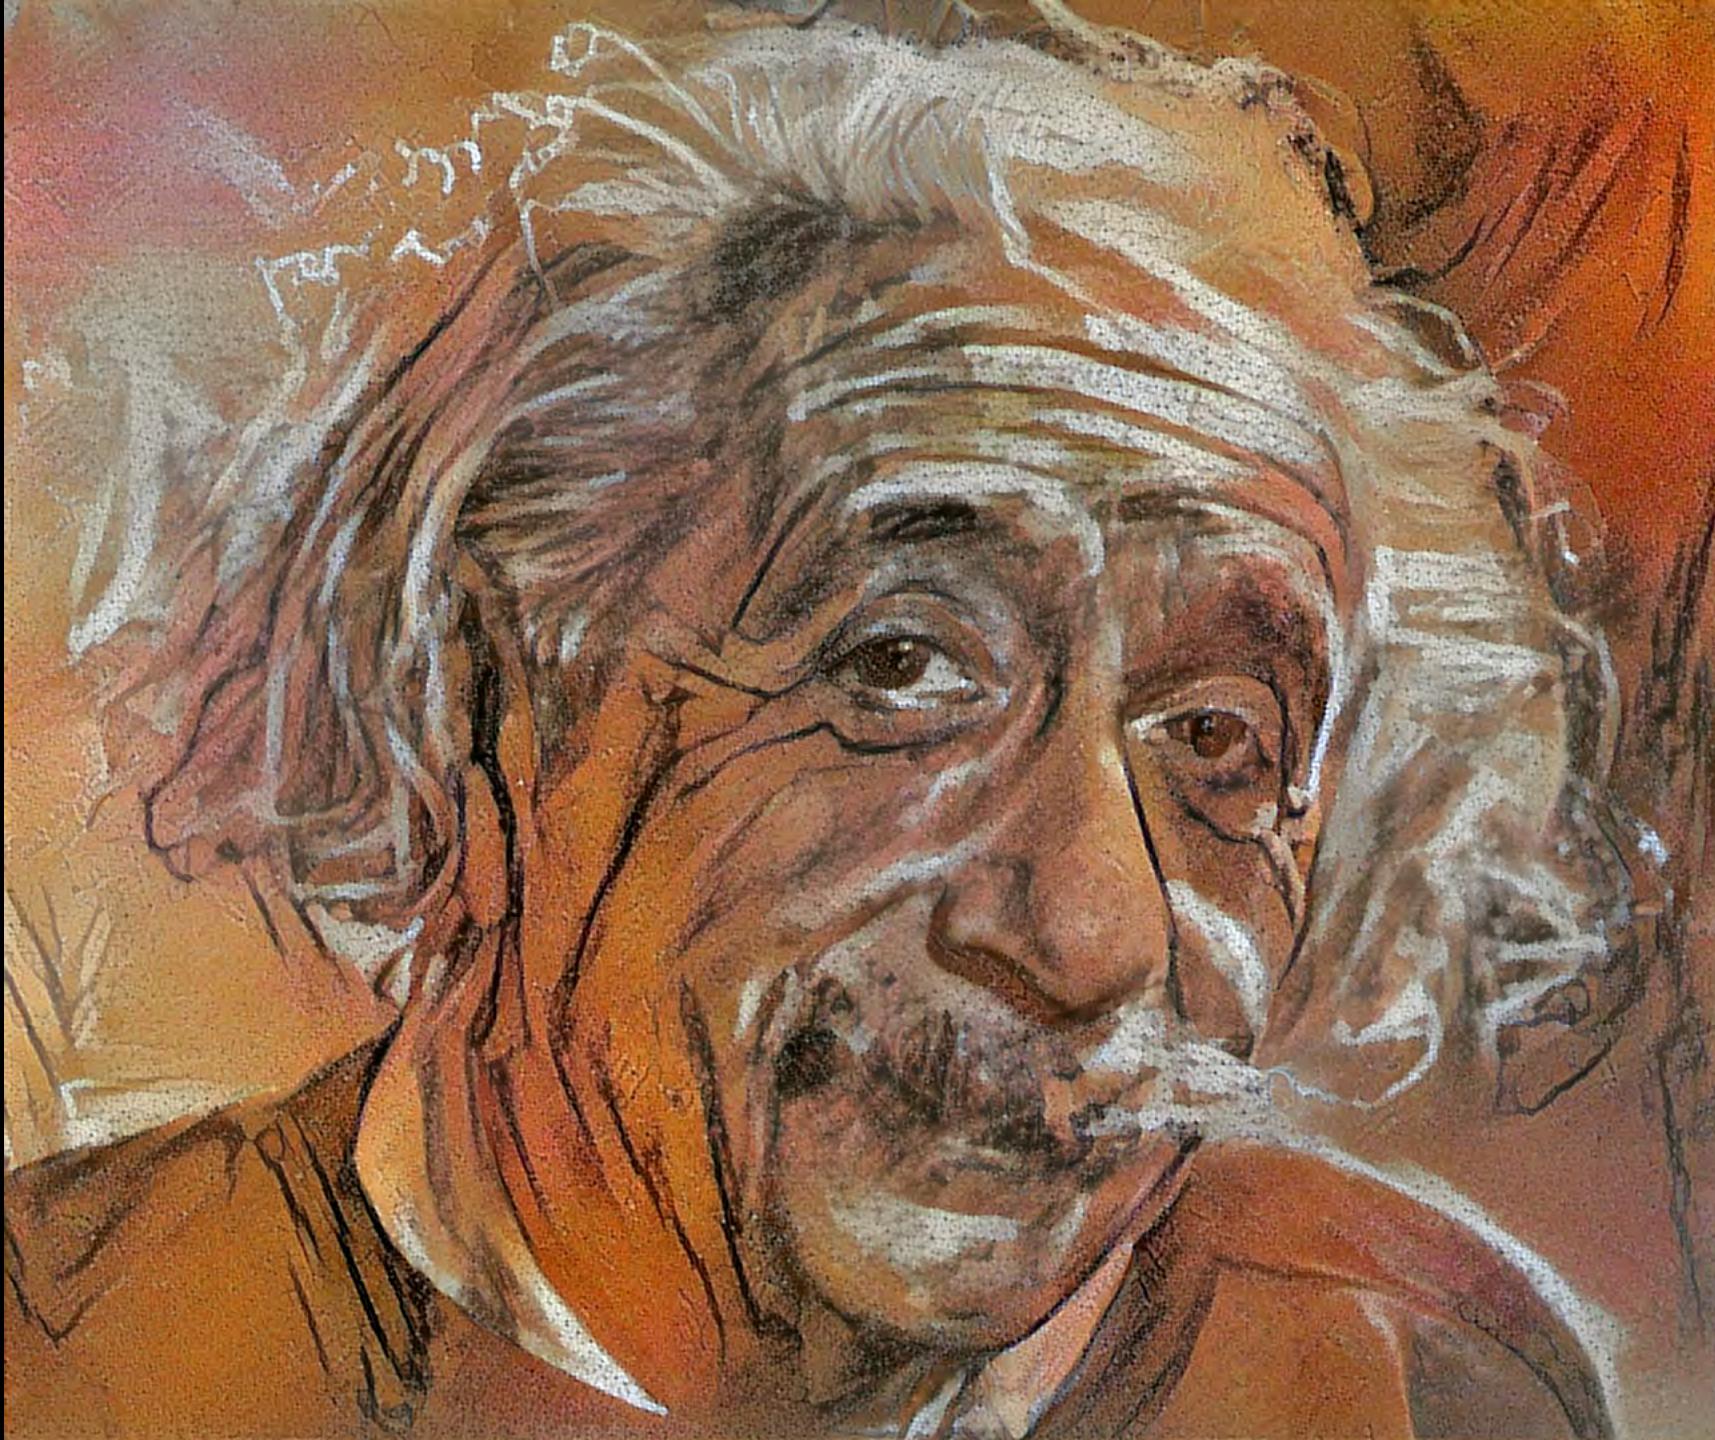
Choose among predefined styles or upload your own style image.



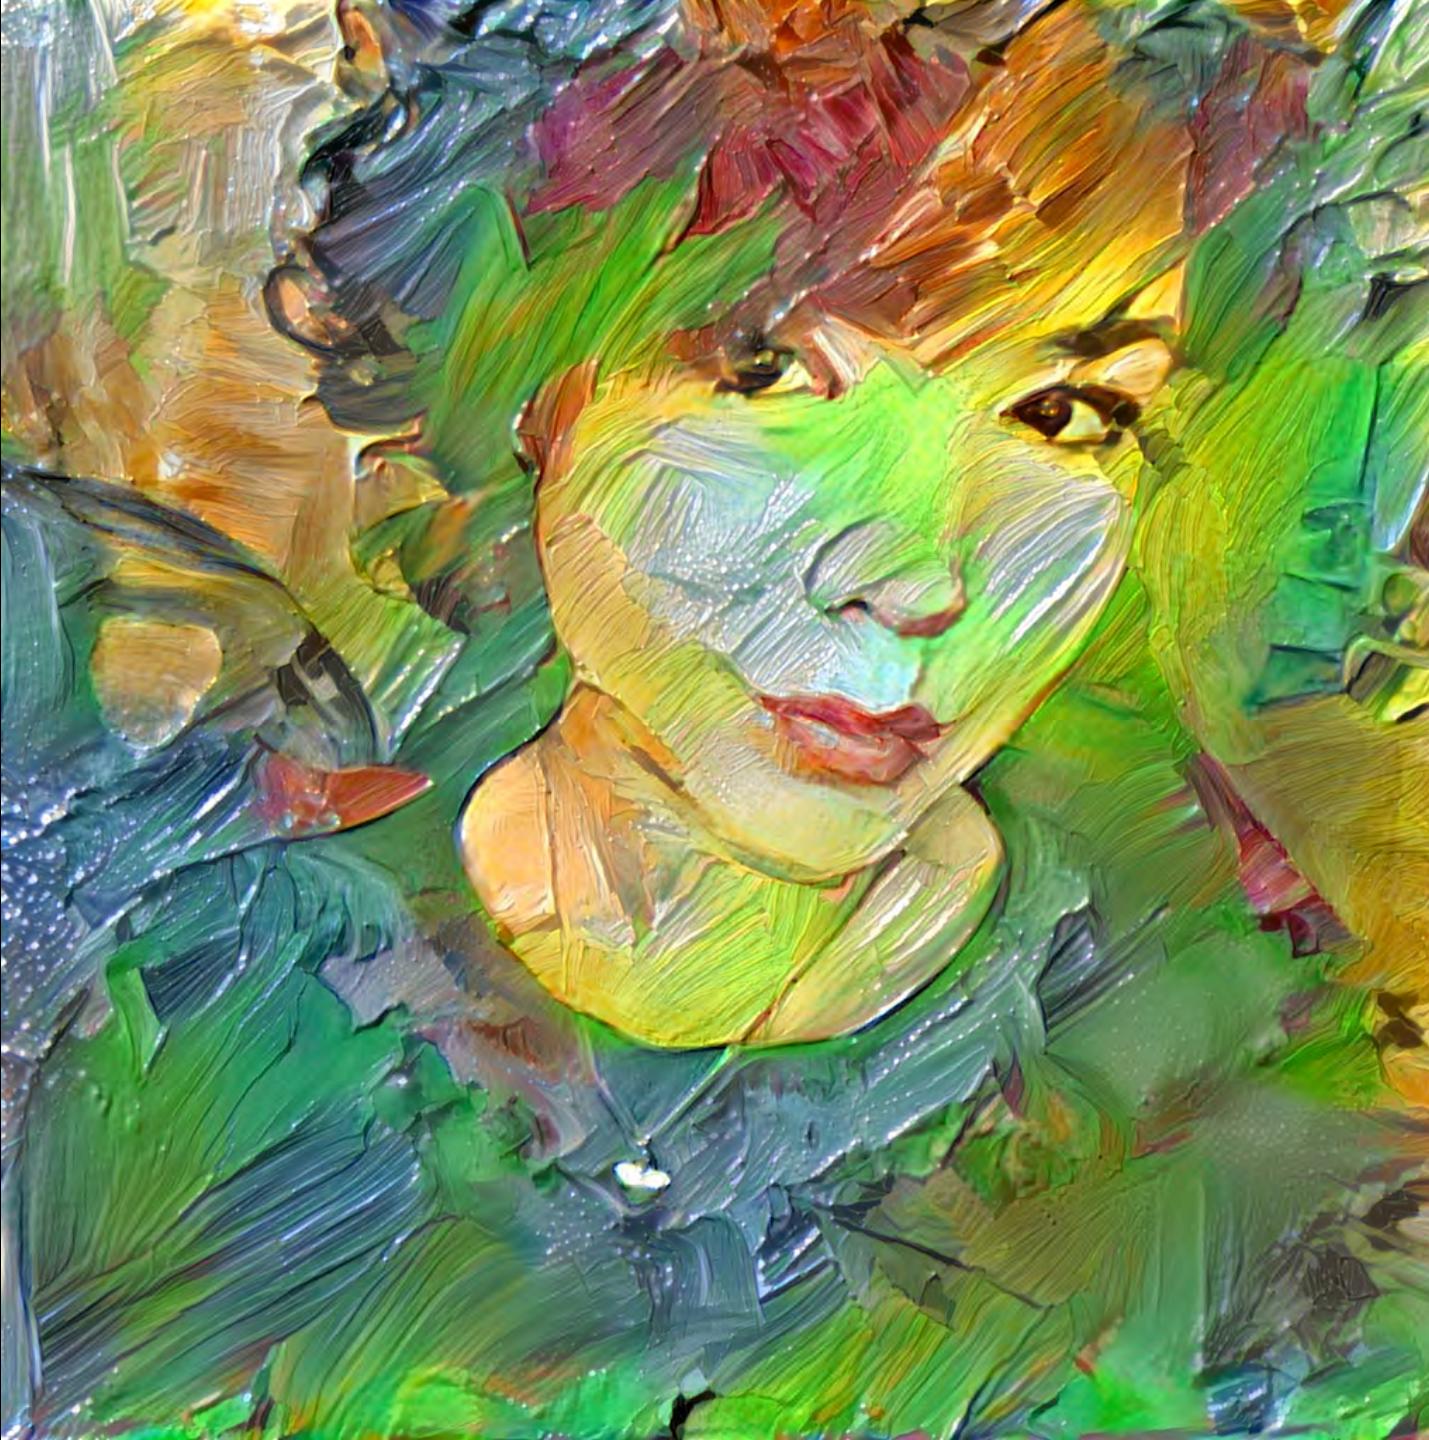
3 Submit

Our servers paint the image for you. You get an email when it's done.













Significance of neural style transfer

Comp. Neuroscience

Texture models
as models of
primate vision

Transfer learning
from deep nets
trained on ImageNet

Machine Learning

Neural Style
Transfer
(Gatys, Ecker, Bethge
CVPR 2016)

Perceptual loss functions
for image processing
(Gatys, Ecker, Bethge CVPR 2017,
Current Opinion Neurobiol. 2017)

Disentangling
content and style
of images

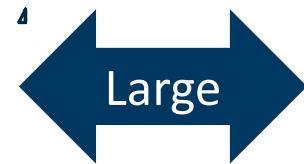
Weaknesses of current
deep learning systems

(Geirhos et al. NeurIPS 2018, Geirhos et
al. ICLR 2019)



Perceptual loss functions

Euclidean distance in pixel space does not cover semantics of images



Perceptual loss functions

Euclidean distance in pixel space does not cover perceptual differences



Blurred



Original



Shifted by 20 pixels



Compute distances in high-level feature space
of deep neural network trained on ImageNet

Neural style transfer implementations

Leon Gatys' PyTorch implementation:

<https://github.com/leongatys/PytorchNeuralStyleTransfer>

Justin Johnson's excellent and very popular Lua/Torch implementation:

(lots of features, but may be hard to run these days)

<https://github.com/jcjohnson/neural-style>

Neural Style in official PyTorch tutorials:

https://pytorch.org/tutorials/advanced/neural_style_tutorial.html

Questions!

-